

Contents

DATA TRANSFER IN ASSEMBLY: WHAT'S REALLY GOING ON	2
DIRECT MEMORY OPERANDS	12
MOV - THE BACKBONE OF ASSEMBLY.....	19
OVERLAPPING REGISTERS AND PARTIAL REGISTER WRITES IN X86	26
SIGNED AND ZERO EXTENSION.....	31
LAHF AND SAHF INSTRUCTIONS	43
DEBUGGING.....	53
ADDITION AND SUBTRACTION	60
CPU STATUS FLAGS/ FLAG REGISTERS: WHY THEY MATTER	70
OFFSET.....	89
HOW OFFSETS WORK IN ASSEMBLY.....	92
ALIGN DIRECTIVE.....	97
TYPE, LENGTHOF, SIZEOF, LABEL	103
DIRECT ADDRESSING MODE.....	112
INDIRECT ADDRESSING MODE	115
INDEXED ADDRESSING	123
IMMEDIATE ADDRESSING MODE	127
STACK ADDRESSING.....	130
POINTERS IN ASSEMBLY LANGUAGE.....	137
TYPEDEF OPERATOR.....	139
JMP AND LOOP INSTRUCTIONS	141
COPYING STRINGS.....	147
CHAPTER SUMMARY: WHAT WE COVERED	160
QUESTIONS.....	161

DATA TRANSFER IN ASSEMBLY: WHAT'S REALLY GOING ON

At its core, **data transfer** in assembly is about moving values around so the CPU can work with them. The processor itself can only operate on data that's in the right place—usually a register—so most assembly programs spend a lot of time copying values between:

- **Registers** (inside the CPU)
- **Memory** (RAM)
- **The stack** (a special area of memory used for function calls and temporary storage)

Nothing “magical” happens during data transfer. The CPU simply copies bits from one location to another. Understanding *where* data comes from and *where* it goes is the foundation of everything else in assembly.

The Three Basic Data Transfer Instructions

MOV

- Copies data from a source to a destination
- Does **not** modify the original source
- Does **not** perform calculations

Think of MOV as a straight copy-paste operation.

PUSH

- Places a value onto the stack
- Automatically adjusts the stack pointer

POP

- Removes a value from the stack
- Stores it somewhere (usually a register)
- Automatically adjusts the stack pointer back

Together, PUSH and POP are essential for function calls, saving registers, and managing temporary data.

Operand Types: What Instructions Work With

Every assembly instruction operates on **operands**. An operand is simply the thing the instruction uses or modifies.

There are **three basic operand types** in x86 assembly:

1. Immediate Operands

Immediate operands are **literal values written directly in the instruction**.

Examples:

- 10
- -255
- 0FFh



Here, 10 is not stored in memory or a register beforehand—it's embedded directly in the instruction.

Why this matters:

Immediate values are fast and convenient, but they're fixed constants. You can't change them at runtime.

2. Register Operands

Register operands refer to **CPU registers**, such as:

- EAX
- EBX
- ECX
- EDX



```
MOV EAX, EBX
```

Registers are:

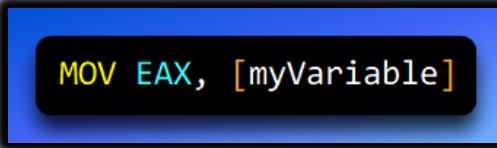
- Extremely fast
- Very limited in number
- Where almost all real computation happens

Key idea:

If the CPU is going to do math or logic, the data usually has to be in registers first.

3. Memory Operands

Memory operands reference **locations in RAM**.



```
MOV EAX, [myVariable]
```

Memory is:

- Much larger than registers
- Slower to access
- Where most program data lives long-term

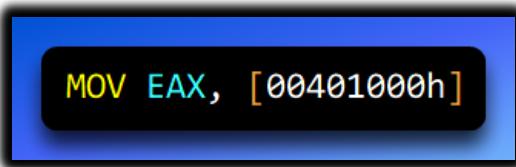
Assembly forces you to be explicit about memory access. You never “accidentally” touch memory—you have to say exactly where.

Addressing Modes: How Memory Is Reached

When an instruction refers to memory, the CPU needs to know **how to find that memory address**. This is where addressing modes come in.

1. Direct Addressing

Direct addressing specifies the memory location explicitly.



What this means:

- myValue represents a fixed memory address
- The CPU goes directly to that address and reads the value

Key characteristics:

- The address is fixed
- Very clear and readable
- Mostly used with labels and global variables

In real programs, direct addressing is common when working with named data defined in the data segment.

2. Immediate Addressing (Not Memory!)

Immediate addressing does **not** access memory at all.



What's happening:

- The value 10 is placed directly into EAX
- No memory lookup occurs

This is included here because beginners often confuse it with memory access — but **it isn't**.

Rule of thumb: No brackets → no memory access

3. Indirect Addressing

Indirect addressing uses a **register that contains a memory address**.



Step-by-step:

1. EBX holds a memory address
2. The CPU looks at that address
3. The value stored there is loaded into EAX

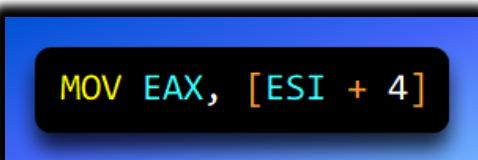
This is **exactly how pointers work** at the assembly level.

Important distinction:

- ebx → the number inside the register
- [ebx] → the data stored at the address contained in EBX

4. Indexed Addressing

Indexed addressing calculates a memory address using a **base register plus an offset**.



What's happening:

- Start at the address in EBX
- Move forward by 4 bytes
- Read the value stored there

This addressing mode is commonly used for:

- Arrays
- Structures
- Walking through memory in loops

Mental model:

“Start here, then move forward this many bytes.”

Indexed addressing is the foundation of data structures in assembly.

Data Transfer Instructions: MOV, PUSH, and POP

With operands and addressing modes understood, data transfer instructions become much clearer.

MOV — Copy Data

MOV copies data from a source to a destination.

Examples:

Operation	Syntax	Example
Register → Register	MOV dest, src	MOV RAX, RBX
Immediate → Register	MOV reg, imm	MOV RAX, 100
Register → Memory	MOV [mem], reg	MOV [var1], RAX
Immediate → Memory	MOV [mem], imm	MOV [var1], 50
Memory → Register	MOV reg, [mem]	MOV RAX, [var1]

⚠️ IMPORTANT RULE

Memory-to-Memory moves are NOT allowed in a single `MOV` instruction. You must move the value to a register first, then to the destination memory address.

Important rule:

- ✖ You **cannot** move memory directly to memory and
- ✓ One operand must be a register.

```
; Invalid  
MOV [var1], [var2]
```

PUSH and POP — Stack Transfers

The stack is a special region of memory managed using the stack pointer (ESP).

```
push eax  
pop ebx
```

What PUSH does:

1. Decreases ESP
2. Stores the value at the new top of the stack

What POP does:

1. Reads the value at the top of the stack
2. Increases ESP

The stack is heavily used for:

- Function calls
- Passing parameters
- Saving registers

Operators That Help with Memory

Assembly provides operators that help calculate and interpret memory addresses.

I. OFFSET

OFFSET gives the **address** of a variable, not its value.

```
mov eax, OFFSET myVar
```

This loads the memory address of myVar into EAX.

II. PTR

PTR tells the assembler **how to treat a memory operand**.

```
mov ax, WORD PTR [ebx]
```

This forces the assembler to treat the memory as a WORD.

This matters because:

- Assemblers do not perform strict type checking
- The CPU needs to know how many bytes to read

III. LENGTHOF

LENGTHOF calculates how many elements are in a data structure.

```
mov ecx, LENGTHOF myArray
```

This is commonly used when writing loops.

Loops and Arithmetic (Preview)

With data transfer understood, you can now:

- Create loops using JMP and LOOP
- Perform arithmetic with ADD, SUB, MUL, and DIV
- Move through arrays and structures using indexed addressing

All of these depend on **correct data movement**.

Flat Memory Model and STDCALL (Windows Context)

When writing 32-bit Windows programs, you'll often see:

```
.MODEL FLAT, STDCALL
```

I. Flat Memory Model

- One continuous 32-bit address space
- No segment juggling
- Memory is treated as a single linear block

This simplifies memory access and matches how modern Windows works.

II. STDCALL Calling Convention

STDCALL defines:

- How parameters are passed (right to left on the stack)
- Who cleans up the stack (the callee)
- How functions interact with the Windows API

This consistency is critical for Windows compatibility.

III. Big Picture Summary

- Data transfer moves values between registers, memory, and the stack
- Operands define *what* data is used
- Addressing modes define *how* memory is reached
- MOV, PUSH, and POP are the core transfer instructions
- OFFSET, PTR, and LENGTHOF help manage memory correctly
- Flat memory and STDCALL define the Windows execution environment

Once this chapter clicks, you're no longer guessing —

DIRECT MEMORY OPERANDS

The real concepts being discussed here are:

1. Direct memory operands (a form of direct addressing)
2. The difference between a value, an address, and the contents at an address
3. How notation (var, [var], hex literals) changes meaning

Direct Memory Operands: Talking to Memory by Name

A **direct memory operand** means:

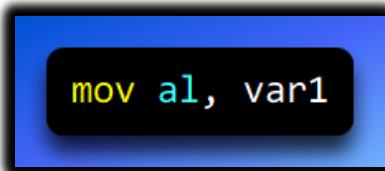
“Access the contents of a specific memory location whose address is known at assembly time.”

In plain English:

- The assembler knows *exactly where this variable lives in memory*
- The instruction hardcodes that address into the machine code

I. What “direct” really means here

When you write:



You are saying:

“Go to the memory location associated with var1, read the byte stored there, and copy it into AL.”

var1 is **not** the value itself.

var1 is a **label** that represents a **fixed memory address**.

II. Declaring a variable: clearing the confusion

```
.data  
var1 BYTE 10h
```

This line means **three separate things**:

1. var1
→ a label (a name for a memory location)
2. BYTE
→ the size of memory being reserved (1 byte)
3. 10h
→ the *initial value* stored in that byte
→ hexadecimal 10h = decimal 16

So what does var1 contain?

- It contains **one byte**
- That byte's value is **16 (decimal)**

What it does *not* mean

- ❌ It does NOT mean “a string”
- ❌ It does NOT mean “hex data is a string”
- ❌ It does NOT mean “10 characters”

Hexadecimal is just a **number format**, not a data type.

III. Hex ≠ string (important mental reset)

This is where many people get tripped up.

- **Hexadecimal** → a way to *write numbers*
- **String** → a sequence of characters stored as numeric codes (ASCII / Unicode)

Example:



Stores:



That byte does *not* represent text unless *you interpret it as text* — and even then, ASCII 16 is a non-printable control character.

So no, a BYTE holding 10h cannot secretly be a string.

IV. Direct memory operand in action

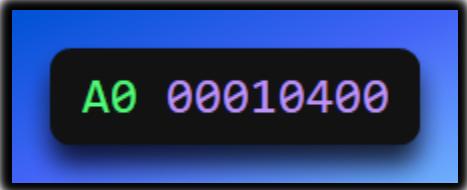


What the CPU actually does:

1. The assembler replaces var1 with its memory address
2. That address is embedded into the instruction
3. At runtime, the CPU:
 - goes to that address
 - reads **1 byte**
 - loads it into AL

V. Why machine code looks like this

We mentioned this:



A0 00010400

Breaking it down:

- A0 → opcode (MOV AL, moffs8)
- 00010400 → 32-bit memory address of var1

This is why it's called **direct**:

the address is literally baked into the instruction.

VI. Direct memory operands with other instructions

Anywhere a memory operand is allowed, direct memory operands can be used:



```
add eax, var1      ; add memory value to register
mov var1, eax      ; store register into memory
```

These all mean:

“Use the contents of the memory location named var1.”

VII. Direct memory + expressions

```
mov al, [var1 + 5]
```

This is called a **direct-offset operand**.

Meaning:

1. Take the address of var1
2. Add 5 bytes
3. Access the memory at that computed address

This is commonly used for:

- arrays
- structure fields
- table lookups

VIII. Why brackets matter (this is HUGE)

No brackets → value or address

```
mov eax, 10000438h
```

This means:

“Load the *number* 10000438h into EAX.”

The CPU does **not** touch memory.

Brackets → dereference (go to memory)

```
mov eax, [10000438h]
```

This means:

“Go to memory address 10000438h and load what’s stored there.”

Golden rule:

Brackets mean “treat this as an address and go there.”

This applies whether the thing inside is:

- a label → [var1]
- a register → [ebx]
- a hex literal → [10000438h]

IX. Can a hex number be a value *or* an address?

Yes — and this is why context matters.

Hex is just a number.

These are different even though the hex is the same:

```
mov eax, 28323428h      ; value
mov eax, [28323428h]    ; memory contents
```

Same number.

Completely different meaning.

X. BYTE vs DWORD (size clarity)

```
var2 BYTE 10h  
var3 DWORD 10248132h
```

- BYTE → 1 byte (0-255)
- DWORD → 4 bytes (32 bits)

The type:

- tells the assembler how much space to reserve
- tells the CPU how many bytes to read/write

It does **not** decide whether something is a value or an address — brackets do.

XI. Opcode vs Operand (clean mental split)

- **Opcode** → what the CPU should do (MOV, ADD, SUB)
- **Operands** → what data the operation works on (registers, memory, immediates)

Example:

```
add eax, 3
```

- Opcode → ADD
- Operands → EAX, 3

XII. Big picture takeaway (this is the anchor)

Direct memory operands mean:

- You name a variable
- The assembler knows its exact address
- The instruction directly accesses that memory location

Brackets decide everything:

- No brackets → value
- Brackets → memory dereference

Once this clicks, pointers, arrays, and structures stop feeling mysterious — they're just **address + size + interpretation**.

MOV - THE BACKBONE OF ASSEMBLY

The MOV instruction as a controlled copy operation, governed by operand roles, sizes, and allowed combinations.

The MOV instruction is used to **copy data from one place to another**.

```
MOV destination, source  
  
; You can also write it in lowercase – assembly is not case-sensitive:  
mov destination, source
```

What MOV actually does

MOV **copies** data.

It does **not**:

- add
- swap
- compare
- or modify the source

After a MOV:

- the **destination changes**
- the **source stays exactly the same**

Think of it like copying text:

You paste the text somewhere else, but the original is untouched.

Destination vs Source (this matters)

```
mov destination, source
```

Destination operand

- The place where the data ends up
- This operand is **modified**

Source operand

- The data being copied
- This operand is **not modified**

Example:

```
mov eax, ebx
```

- EAX → destination (changed)
- EBX → source (unchanged)

Operand sizes must match

Both operands must be the **same size**. Valid:

```
mov al, bl          ; 8-bit to 8-bit  
mov ax, bx          ; 16-bit to 16-bit  
mov eax, ebx        ; 32-bit to 32-bit
```

Invalid:

```
mov eax, al          ; size mismatch
```

The CPU needs to know *exactly how many bytes* to move.

MOV cannot move memory to memory (and why)

This is one of the most important rules:

✖ **MOV cannot copy data directly from one memory location to another.**

This is not a syntax limitation — it's a **CPU rule**.

Memory-to-memory moves would require:

- two memory reads
- one memory write
- extra internal buffering

x86 keeps things simple and fast by requiring **at least one operand to be a register**.

How to move data from memory to memory (the correct way)

You use a register as a temporary step.

```
.data  
var1 WORD ?  
var2 WORD ?  
  
.code  
mov ax, var1      ; memory → register  
mov var2, ax      ; register → memory
```

What's happening:

1. The value of var1 is loaded into AX
2. The value in AX is stored into var2

Result:

- var2 now contains the same value as var1

Register → Memory example

```
.data  
var3 DWORD ?  
  
.code  
mov eax, 1024  
mov var3, eax
```

Step-by-step:

1. 1024 is loaded into EAX
2. The contents of EAX are copied into memory at var3

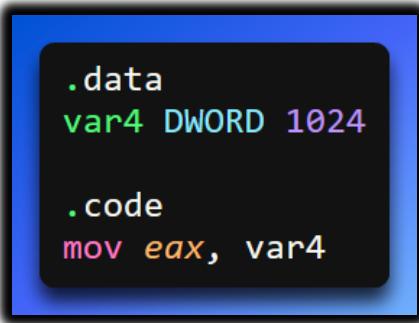
Important idea:

- var3 is just a **name for a memory address**
- The assembler reserves **4 bytes** because it is a DWORD

After execution:

- var3 holds the value 1024

Memory → Register example



```
.data  
var4 DWORD 1024  
  
.code  
mov eax, var4
```

What happens:

1. The CPU goes to the memory location named var4
2. Reads 4 bytes
3. Copies them into EAX

After this instruction:

- EAX = 1024

Immediate values with MOV

You can also move **immediate (literal) values** into registers or memory.

```
mov eax, 5  
mov var3, 20
```

Rules:

- Immediate → register ✓
- Immediate → memory ✓
- Immediate → immediate ✗ (doesn't make sense)

Summary: All valid MOV forms

```
mov reg, reg      ; register to register  
mov reg, mem      ; memory to register  
mov mem, reg      ; register to memory  
mov reg, imm      ; immediate to register  
mov mem, imm      ; immediate to memory
```

And the one that's **not allowed**:

```
mov mem, mem      ; ✗ invalid
```

Why MOV is so important

Almost everything in assembly depends on MOV:

- loading values
- storing results
- passing function parameters
- working with memory
- preparing data for arithmetic

If you understand:

- **destination vs source**
- **size matching**
- **register involvement**

then you understand how data flows through the CPU.

Big picture takeaway

MOV is not “move” — it’s **copy**.

It copies:

- values
- addresses
- memory contents

But it always follows strict rules so the CPU knows:

- how much data to copy
- where it’s coming from
- where it’s going

Master MOV, and assembly stops feeling chaotic — it becomes **deliberate and predictable**.

OVERLAPPING REGISTERS AND PARTIAL REGISTER WRITES IN X86

More specifically:

- How **AL, AX, and EAX share the same physical storage**
 - How **writing to a smaller part of a register affects (or does not affect) the rest**
-

Overlapping Values: How Partial Register Writes Work

In x86 assembly, many registers **overlap**.

This means that smaller registers are **not separate storage** — they are **views into a larger register**.

The best example is EAX.

I. The EAX register layout (mental model)

A 32-bit register like EAX is divided like this:



- **AL** → lowest 8 bits
- **AH** → next 8 bits
- **AX** → lowest 16 bits (AH + AL)
- **EAX** → all 32 bits

They all refer to the **same physical register**.

II. Data declarations (what we're loading)

```
.data  
oneByte    BYTE    78h  
oneWord    WORD    1234h  
oneDword   DWORD   12345678h
```

- oneByte → 8 bits → 78h
- oneWord → 16 bits → 1234h
- oneDword → 32 bits → 12345678h

Step-by-step: how EAX changes

I. Clear EAX

```
mov eax, 0
```

All 32 bits are zeroed: **EAX = 00000000h**

II. Move a BYTE into AL

```
mov al, oneByte
```

What happens:

- Only **AL** (lowest 8 bits) is overwritten
- The upper 24 bits remain unchanged

EAX = 00000078h

Key idea: Writing to AL affects only 1 byte.

III. Move a WORD into AX

```
mov ax, oneWord
```

What happens:

- The lower **16 bits** are overwritten
- The upper **16 bits** remain unchanged

EAX = 00001234h

Notice:

- The previous 78h in AL is gone
- Because AL is part of AX

IV. Move a DWORD into EAX

```
mov eax, oneDword
```

What happens:

All 32 bits are replaced - **EAX = 12345678h**

No overlap concerns here — the entire register is rewritten.

V. Zero AX only

```
mov ax, 0
```

What happens:

- Lower 16 bits → set to zero
- Upper 16 bits → **unchanged**

EAX = 12340000h

This is where people get surprised.

VI. Why this matters (the dangerous part)

Because partial writes **do not clear the rest of the register**, you can easily end up with **garbage in the upper bits** if you're not careful.

Example bug:

```
mov al, 1  
; programmer assumes EAX = 1
```

Reality: EAX = ??????01h

Unless EAX was cleared earlier, the upper bits contain whatever was there before.

VII. Golden rules for overlapping registers

Rule 1: Smaller writes do not clear larger registers

- mov al, x → changes 8 bits
- mov ax, x → changes 16 bits
- mov eax, x → changes all 32 bits

Rule 2: If you care about the full register, write the full register

Instead of:

```
mov al, value  
; safer  
  
xor eax, eax  
mov al, value
```

Now you *know* the upper bits are zero.

Rule 3: Overlap works both ways

- Writing to AX overwrites AL
- Writing to AL does **not** preserve meaningful values in AX
- Everything overlaps downward

VII. Why x86 works this way

This design exists for:

- backward compatibility (8-bit and 16-bit CPUs)
- performance
- flexibility

It's powerful — but sharp.

VIII. Big picture takeaway

AL, AX, and EAX are not separate registers.

They are different-sized windows into the same storage.

When you move smaller data into a larger register:

- it goes into the **lower portion**
- the **upper bits are untouched**

Understanding this prevents:

- subtle bugs
- incorrect comparisons
- broken arithmetic
- mysterious values

Once this clicks, you start *controlling* the CPU.

SIGNED AND ZERO EXTENSION

Table 4-1 Instruction Operand Notation, 32-Bit Mode.

Operand	Description
<i>reg8</i>	8-bit general-purpose register: AH, AL, BH, BL, CH, CL, DH, DL
<i>reg16</i>	16-bit general-purpose register: AX, BX, CX, DX, SI, DI, SP, BP
<i>reg32</i>	32-bit general-purpose register: EAX, EBX, ECX, EDX, ESI, EDI, ESP, EBP
<i>reg</i>	Any general-purpose register
<i>sreg</i>	16-bit segment register: CS, DS, SS, ES, FS, GS
<i>imm</i>	8-, 16-, or 32-bit immediate value
<i>imm8</i>	8-bit immediate byte value
<i>imm16</i>	16-bit immediate word value
<i>imm32</i>	32-bit immediate doubleword value
<i>reg/mem8</i>	8-bit operand, which can be an 8-bit general register or memory byte
<i>reg/mem16</i>	16-bit operand, which can be a 16-bit general register or memory word
<i>reg/mem32</i>	32-bit operand, which can be a 32-bit general register or memory doubleword
<i>mem</i>	An 8-, 16-, or 32-bit memory operand

Zero Extension (moving a smaller unsigned value into a larger register)

Zero extension is used when you move a **smaller unsigned value** (for example, 8-bit or 16-bit) into a **larger register** (like a 32-bit register), and you want the extra space to be filled with **zeros**.

This is important because unsigned values do **not** have a sign. So, when we make them bigger, we must not copy or invent a sign bit. We simply add zeros to the left.

The basic idea is simple:

1. First, clear the larger register (set it to zero).
2. Then, move the smaller value into the lower part of that register.
3. The upper bits stay zero, which is exactly what we want.

Example setup

Let's say we have a variable called count:

- count is a **16-bit unsigned integer**
- Its value is **1**
- In hexadecimal, that value is **0001h**

Assembly code example (Zero Extension)

```
mov ecx, 0          ; Clear the entire 32-bit ECX register
mov cx, count       ; Move the 16-bit value into the lower 16 bits of ECX
```

What happens step by step

Step 1: Clear ECX

```
mov ecx, 0
```

This sets all 32 bits of the ECX register to zero: **ECX = 00000000h**

This step is important because it guarantees that the upper 16 bits are zero before we move anything into ECX.

Step 2: Move the 16-bit value into CX

```
mov cx, count
```

- CX is the **lower 16 bits** of ECX
- The value of count is 0001h
- This instruction copies 0001h into the lower 16 bits only

After this instruction:

- Upper 16 bits of ECX: 0000h
- Lower 16 bits of ECX: 0001h

So, the full 32-bit value of ECX becomes: **ECX = 00000001h**

We took a **16-bit unsigned value** and placed it into a **32-bit register**.

The extra 16 bits on the left were filled with **zeros**, not sign bits.

That is why this is called **zero extension**.

Final result (important correction)

- Hex value in ECX: 00000001h
- Decimal value in ECX: **1**

We copied the number **1** into ECX and padded it on the left with zeros.
The value stays **1**, not 16.

MOVZX and MOVSX Instructions

(Move with Zero Extension / Move with Sign Extension)

When moving **smaller values** (8-bit or 16-bit) into **larger registers** (16-bit or 32-bit), we often need to **extend** the value.

Intel provides two special instructions that do this automatically:

- **MOVZX** → for **unsigned values**
- **MOVSX** → for **signed values**

These instructions save us from manually clearing registers and make our code **shorter, safer, and easier to read**.

Instruction Forms

I. MOVZX (Zero Extension)

```
MOVZX reg32, reg/mem8  
MOVZX reg32, reg/mem16  
MOVZX reg16, reg/mem8
```

II. MOVSX (Sign Extension)

```
MOVSX reg32, reg/mem8  
MOVSX reg32, reg/mem16  
MOVSX reg16, reg/mem8
```

MOVZX — Move with Zero Extension

MOVZX copies a smaller unsigned value into a larger register and fills all the extra bits with zeros.

It does **not care about signs**.

It treats the value as always positive.

I. Example: Zero extension with MOVZX

Data section

```
.data  
count WORD 1      ; 16-bit unsigned integer (0001h)
```

Code section

```
.code  
movzx ecx, count
```

What happens step by step

- count is **16 bits** and equals 0001h
- ECX is **32 bits**
- MOVZX:
 - Copies count into the lower 16 bits of ECX
 - Fills the upper 16 bits with **zeros**

Final result: **ECX = 00000001h**

Decimal value: **1**

This is **zero extension** — the value is extended by adding zeros on the left.

MOVSX — Move with Sign Extension

MOVSX copies a smaller signed value into a larger register and preserves the sign.

- If the value is **positive**, it fills the upper bits with **0**
- If the value is **negative**, it fills the upper bits with **1**

This keeps the number mathematically correct after the move.

I. Example: Negative signed value

Data section and code section

```
.data  
signedVal SWORD -16    ; 16-bit signed integer (FFF0h)  
  
; code section  
  
.code  
movsx ecx, signedVal
```

What happens step by step

- signedVal = -16
- Hex value (16-bit): FFF0h
- Sign bit = 1 (negative)

After MOVSX: **ECX = FFFFFFFF0h**

The upper bits are filled with **1s**, preserving the negative sign.

Important Question:

Do signed values have to be negative for MOVSX?

No. Absolutely not.

MOVSX works for **both positive and negative signed values**.

Example: Positive signed value

Data section and Code section

```
.data  
signedVal SWORD 42      ; 16-bit signed integer (002Ah)  
  
; code section  
  
.code  
movsx ecx, signedVal
```

What happens here?

- 42 in hex is 002Ah
- Sign bit = 0 (positive)

After MOVSX: ECX = **0000002Ah**

The value stays positive, and the upper bits are filled with zeros.

✓ **Sign preserved correctly**

Key Rule for Sign Extension (Very Important)

When using **MOVSX**, always look at the **most significant bit (sign bit)**:

- **Sign bit = 0** → positive → extend with **zeros**
- **Sign bit = 1** → negative → extend with **ones**

Decimal	Hex Value	Sign Bit	Extended With
42	002Ah	0	Zeros (Positive)
-6	FFF4h	1	Ones (Negative)
-16	FFF0h	1	Ones (Negative)

Why do we need MOVZX if MOVSX exists?

Great question — and this is where many learners get confused.

Short answer: Because MOVSX can break unsigned values.

MOVZX — for unsigned data

- Always fills upper bits with **zeros**
- Safe for values like counters, sizes, indexes

Example:

```
movzx ecx, count
```

MOVSX — for signed data

- Copies the **sign bit**
- Correct for negative numbers
- Dangerous for unsigned values

Example problem:

```
count WORD 0FFEh    ; unsigned = 65534
movsx ecx, count    ; WRONG!

; Result = ECX = FFFFFFFEh - Interpreted as -2
```

⚠ That's why **MOVZX** exists.

Final Summary (Simple and Clear)

- **MOVZX**
 - Used for **unsigned values**
 - Upper bits are filled with **zeros**
 - Prevents accidental negative values
- **MOVSX**
 - Used for **signed values**
 - Upper bits copy the **sign bit**
 - Preserves positive or negative meaning

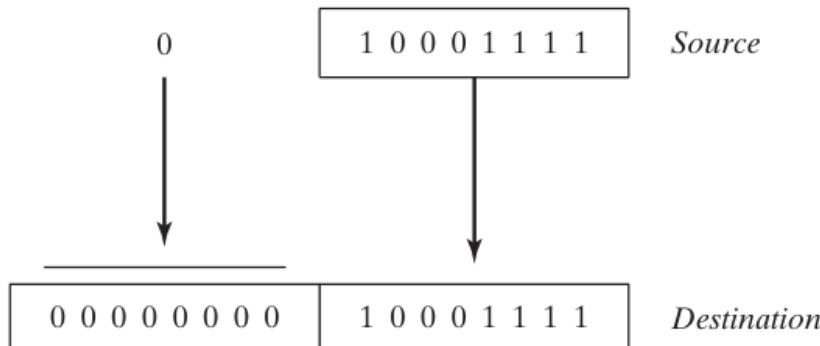
👉 Always choose based on the data type, not the instruction size.

Movzx Repeated

```
.data  
    byteVal BYTE 10001111b      ;decimal 143  
  
.code  
    movzx ax,byteVal           ; AX = 0000000010001111b
```

What is happening here:

FIGURE 4–1 Using MOVZX to copy a byte into a 16-bit destination.



MOVZX, which stands for Move with Zero Extension, is an instruction that lets us take a value that is stored in a smaller register or memory location and copy it into a bigger register.

When we do this, the extra space in the bigger register (the upper bits) is filled with zeros so that the value stays the same. Let's look at some examples using registers to understand how this works step by step.

1. Example: Register Operands

```
mov bx, 0A69Bh      ; BX = 0A69Bh
movzx eax, bx        ; EAX = 0000A69Bh
movzx edx, bl        ; EDX = 0000009Bh
movzx cx, bl         ; CX = 009Bh
```

In this example, we used the **MOVZX** instruction to copy smaller pieces of data from one place into bigger registers, making sure that any extra space is filled with zeros. Let's take it one step at a time:

1. **mov bx, 0A69Bh** – Here, we put the 16-bit value 0A69Bh into the BX register. Think of BX as a small container that can hold 16 bits (2 bytes).
2. **movzx eax, bx** – Now, we want to copy that 16-bit value from BX into the larger 32-bit register EAX. Because EAX is bigger, MOVZX fills the extra 16 bits at the top with zeros. After this, EAX contains 0000A69Bh.
3. **movzx edx, bl** – BL is the lower 8 bits of BX (just the last byte). We copy this into the 32-bit register EDX using MOVZX. The top 24 bits of EDX are filled with zeros, so the final value is 0000009Bh.
4. **movzx cx, bl** – Finally, we copy that same 8-bit value from BL into the 16-bit register CX. Since CX has space for 16 bits, the upper 8 bits are filled with zeros, giving 009Bh.

2. Example: Memory Operands

```
.data
    byte1 BYTE 9Bh
    word1 WORD 0A69Bh

.code
    movzx eax, word1      ; EAX = 0000A69Bh
    movzx edx, byte1      ; EDX = 0000009Bh
    movzx cx, byte1       ; CX = 009Bh
```

In this example, we used **MOVZX** again, but this time the data comes from memory instead of registers. We have two variables in memory: byte1 (8 bits) and word1 (16 bits). Here's what happens step by step:

1. **movzx eax, word1** – We take the 16-bit value stored in word1 and copy it into the 32-bit register EAX. Since EAX is bigger than word1, MOVZX fills the extra 16 bits at the top with zeros. After this, EAX contains 0000A69Bh.
2. **movzx edx, byte1** – Now we take the 8-bit value in byte1 and copy it into the 32-bit register EDX. The top 24 bits are filled with zeros, so EDX becomes 0000009Bh.
3. **movzx cx, byte1** – Finally, we copy the same 8-bit value from byte1 into the 16-bit register CX. The upper 8 bits of CX are filled with zeros, giving 009Bh.

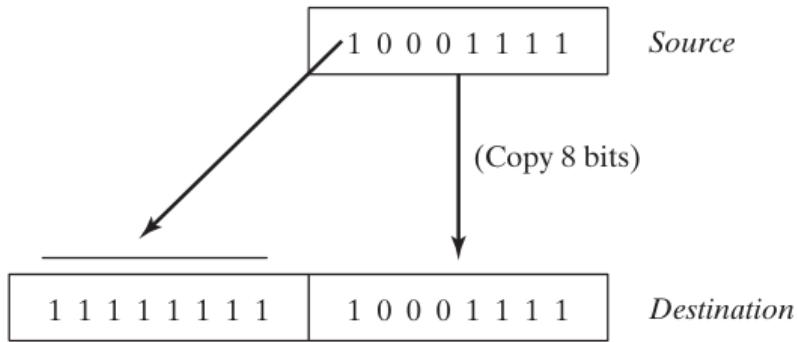
The important idea here is that **MOVZX always fills the extra bits with zeros** when moving smaller values into larger registers. This keeps the value correct and avoids accidentally changing the sign of the number.

MOVSX Repeated

```
;MOVSX  
  
mov bx, 0A69Bh      ; BX = 0A69Bh  
movsx eax, bx        ; EAX = FFFFA69Bh  
movsx edx, bl        ; EDX = FFFFFFF9Bh  
movsx cx, bl          ; CX = FF9Bh
```

What is happening here?

FIGURE 4–2 Using MOVSX to copy a byte into a 16-bit destination.



Here, we are looking at **MOVSX**, which copies smaller values into larger registers while **preserving the sign** (positive or negative). We're using the hexadecimal value 0A69Bh as an example. Here's what happens step by step:

1. **mov bx, 0A69Bh** – This puts the 16-bit value 0A69Bh into the BX register. The first digit "A" in hexadecimal means that the highest bit (the **sign bit**) is 1, so this value is considered **negative** in signed representation.
2. **movsx eax, bx** – We copy the 16-bit value in BX into the 32-bit register EAX. Because we are using MOVSX, it **keeps the sign the same**. The upper 16 bits of EAX are filled with ones (F in hex) to preserve the negative sign. After this, EAX = FFFFA69Bh.
3. **movsx edx, bl** – Now we take the lower 8 bits of BX (called BL) and copy them into the 32-bit EDX register. MOVSX again preserves the sign. The upper 24 bits of EDX are filled with ones, so EDX = FFFFFFF9Bh.
4. **movsx cx, bl** – Finally, we copy the 8-bit BL into the 16-bit register CX. The upper 8 bits of CX are filled with ones because the sign bit is 1, giving CX = FF9Bh.

Key point: The "A" in 0A69Bh shows that the number is negative in signed form. MOVSX makes sure that when we copy smaller numbers into larger registers, the **sign is preserved** by filling the extra bits with ones for negative numbers.

LAHF AND SAHF INSTRUCTIONS

The **LAHF** (Load AH from Flags) and **SAHF** (Store AH into Flags) instructions are used to **move specific CPU flags to and from the AH register**.

These instructions are useful when you want to **save or restore the status of certain flags** without affecting the rest of the EFLAGS register.

LAHF — Load AH from Flags

I. What it does

- The **LAHF** instruction takes the **low byte of the FLAGS register** (specifically the SF, ZF, AF, PF, and CF flags) and **copies it into the AH register**.
- This allows you to **save the state of these flags** for later use, such as storing them in memory or comparing them without changing the CPU state.

II. Flags included in LAHF

Flag	Meaning	Condition for Set (1)
SF	Sign Flag	The Most Significant Bit (MSB) of the result is 1 (Negative).
ZF	Zero Flag	All bits of the result are 0.
AF	Auxiliary Flag	Carry/Borrow out of bit 3 (Used for BCD arithmetic).
PF	Parity Flag	The lowest byte of the result contains an even number of 1s.
CF	Carry Flag	An unsigned overflow occurred (carry out or borrow in).

Note: LAHF **does not load all EFLAGS**, only these selected flags.

III. Example: Using LAHF

```
; Save the state of SF, ZF, AF, PF, and CF into a variable  
lahf          ; Load flags into AH  
mov [savedFlags], ah ; Store AH into memory
```

Or

```
.data  
    saverflags BYTE ?    ;Define a variable to store the flags  
  
.code  
    lahf          ;load flags into AH  
    mov saverflags, ah ;save them in the 'saverflag' variable
```

Step-by-step explanation:

1. lahf copies the **low byte of FLAGS** into the AH register.
2. mov [savedFlags], ah stores the value of AH into a memory variable called savedFlags.
3. Later, you can **restore these flags** using SAHF.

In this example, after executing LAHF, the AH register contains the values of the specified flags, and you can save these values in the saveflags variable for future reference.

SAHF (Store AH into Flags):

The SAHF instruction works in the opposite direction. It copies the value from the AH register into the low byte of the EFLAGS (or RFLAGS) register.

This allows you to restore saved flag values.

Here's an example of how to use SAHF to retrieve saved flag values from a variable:

```
.data  
    saveflags BYTE ? ; Variable containing saved flags  
  
.code  
    mov ah, saveflags ; Load saved flags into AH  
    sahf              ; Copy AH into the Flags register
```

EFLAGS / RFLAGS Register

The **EFLAGS** register (called **RFLAGS** in 64-bit mode) is like a **control panel for the CPU**.

- It's **32 bits (4 bytes)** in 32-bit mode, and **64 bits (8 bytes)** in 64-bit mode.
- It **stores status flags** that tell the CPU what happened after each instruction.
- It also has **control bits** that affect how the CPU works.

Some of the important **flags** you'll see in EFLAGS/RFLAGS are:

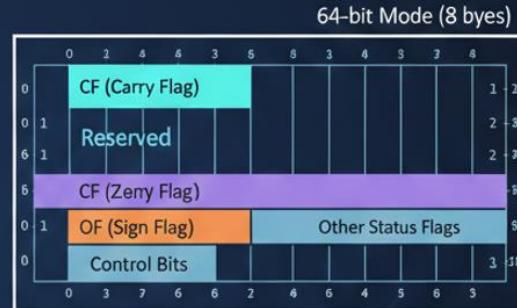
- **Carry Flag (CF)**: Tells if a calculation carried out of the highest bit.
- **Zero Flag (ZF)**: Tells if the result of a calculation was zero.
- **Sign Flag (SF)**: Tells if the result was negative.
- **Overflow Flag (OF)**: Tells if a calculation overflowed the allowed range.
- ...and many others that help the CPU **track what's going on**.

In **64-bit mode**, the **RFLAGS register** works the same way as EFLAGS, just **bigger (64 bits)**.

32-bit Mode (EFLAGS)



64-bit Mode (RFLAGS)



- ⊕ Carry Flag (CF): Carried out of the highest bit.
- ✓ Sign Flag (ZF): Result was negative.
- ⊖ Overflow Flag (OF): Result exceeded range.

XCHG (Exchange) Instruction

The XCHG instruction in x86 assembly is used to **swap the contents of two operands**.

Unlike doing a manual swap using a temporary register, XCHG can handle this directly, which makes your code shorter, cleaner, and sometimes faster.

There are **three main variants**:

1. XCHG reg, reg – exchanges contents between two registers.
2. XCHG reg, mem – exchanges contents between a register and a memory location.
3. XCHG mem, reg – exchanges contents between a memory location and a register.

Let's break each one down, with examples and practical notes.

1. XCHG reg, reg

What it does: Swaps the contents of two registers.

- No memory access is required—everything stays inside the CPU registers.
- Useful when you need to reorder values quickly, e.g., swapping two counters.

Example:

```
MOV AX, 5      ; AX = 5
MOV BX, 10     ; BX = 10
XCHG AX, BX   ; Swap AX and BX
```

Result after execution:

- AX = 10
- BX = 5

Key Points:

- Fast, because it only involves registers.
- Atomic on most processors, which means it can safely be used in some synchronization tasks (like implementing locks).

2. XCHG reg, mem

What it does: Swaps the value of a register with a value in memory.

- Now the CPU has to read from and write to RAM, which is slower than register-only operations.
- This is very handy when you need to update a variable without using an extra temporary register.

Example:

```
MOV BX, 20      ; BX = 20
XCHG BX, [var1] ; Swap BX with memory location 'var1'
```

Result:

- BX now contains the previous value of var1.
- var1 now contains 20.

Why it matters:

- Helps in **swapping array elements** directly in memory.
 - Can also be used in simple **locking mechanisms**, because XCHG is atomic when working with a register and memory.
-

3. XCHG mem, reg

What it does: Swaps a memory location with a register. This is technically the same as XCHG reg, mem because the CPU treats memory-register swaps identically, but conceptually it's about **memory being the first operand**.

Example (swapping two memory locations using a temporary register):

```
MOV AX, [val1]    ; AX = val1
XCHG AX, [val2]    ; Swap AX with val2
MOV [val1], AX    ; Move AX (original val2) back to val1
```

Result:

- val1 and val2 are swapped.
- AX temporarily holds the value during the swap.

Why it matters:

- Direct memory-to-memory swap isn't supported on x86, so this pattern (using a register as a "bridge") is essential.
- Shows how XCHG can be combined with MOV to manipulate memory efficiently.

Practical Tips & Takeaways

- **Atomicity:** XCHG is inherently atomic when one operand is a memory location. This makes it useful in multithreading scenarios.
- **Performance:** Register-register swaps are fastest; memory involvement slows things down.
- **Common Uses:**
 - Swapping variables (obvious!)
 - Sorting algorithms (like bubble sort)
 - Implementing simple locks or semaphores in low-level concurrent code

Direct-Offset Operands

Direct-offset operands let you **access memory at a specific location** by starting from a variable's base address and adding an **offset**.

```
arrayB BYTE 10h, 20h, 30h, 40h, 50h
mov al, arrayB           ;al = 10h
mov al, [arrayB+1]       ;al = 20h
mov al, [arrayB+2]       ;al = 30h
```

- They're useful for **arrays or structures**, where you want to pick a specific element.
- The offset is just a number that tells the CPU how far from the base address to go.

Example concept:

- If you have arrayB, then arrayB + 20 is the **effective address** of the second element (first element is at arrayB + 0).
- Adding brackets, like [arrayB + 20], tells the assembler: "I want the **value stored at this memory location**, not the address itself."

```
mov al, [arrayB+20]      ;al = retrieving memory outside the array
```

Tip:

- You don't **have to** use brackets in MASM, but it's **much clearer if you do**.
- MASM **doesn't check if your address is in range**, so if your array has 20 bytes and you access arrayB + 20, you'll read **memory outside the array**. This won't cause an immediate error but can create a **sneaky logic bug**.

I. Word and Doubleword Arrays

When accessing word and doubleword arrays, you need to take into account the size of the array elements.

For example, the following code shows how to access the second element in an array of 16-bit words:

```
.data  
    arrayW WORD 100h,200h,300h  
.code  
    mov ax,[arrayW+2] ; AX = 200h
```

To put it simply, **[arrayW + 2]** is a way of telling the computer: "Start at the beginning of the list called **arrayW**, skip the first item, and look at the second one."

II. Why do we add 2?

Computers don't see "items" in a list; they see **bytes** (tiny units of storage).

- In this specific list (**arrayW**), every single item takes up **2 bytes** of space.
- To get to the second item, you have to jump over those first 2 bytes.

The code **[arrayW+2]** tells the computer to skip the first 2 bytes of the list. Since each item in this list is 2 bytes big, this jump moves you exactly to the start of the second item.

In the same way, if the items are "doublewords" (4 bytes big), you would add 4 to reach the second item.

```
.data  
    arrayD DWORD 10000h,20000h  
.code  
    mov eax,[arrayD+4] ; EAX = 20000h
```

The code **[arrayD+4]** tells the computer to skip the first 4 bytes to reach the second item in the list. This is because each item in this list is 4 bytes wide.

You must be careful: the computer will not stop you if you jump too far.

If you use a number that is too large, you might accidentally look at memory that doesn't belong to your program, which can cause errors.

```
.386
.model flat, stdcall
.stack 4096

ExitProcess PROTO, dwExitCode:DWORD

.data
    val1 WORD 1000h
    val2 WORD 2000h
    arrayB BYTE 10h,20h,30h,40h,50h
    arrayW WORD 100h,200h,300h
    arrayD DWORD 10000h,20000h

.code
main PROC
; Demonstrating MOVZX instruction:
    mov bx, 0A69Bh
    movzx eax, bx ; EAX = 0000A69Bh
    movzx edx, bl ; EDX = 0000009Bh
    movzx cx, bl ; CX = 009Bh

; Demonstrating MOVSX instruction:
    mov bx, 0A69Bh
    movsx eax, bx ; EAX = FFFFA69Bh
    movsx edx, bl ; EDX = FFFFFFF9Bh
    mov bl, 7Bh
    movsx cx, bl ; CX = 007Bh

; Memory-to-memory exchange:
    mov ax, val1 ; AX = 1000h
    xchg ax, val2 ; AX=2000h, val2=1000h
    mov val1, ax ; val1 = 2000h

; Direct-Offset Addressing (byte array):
    mov al, arrayB ; AL = 10h
    mov al, [arrayB+1] ; AL = 20h
    mov al, [arrayB+2] ; AL = 30h

; Direct-Offset Addressing (word array):
    mov ax, arrayW ; AX = 100h
    mov ax, [arrayW+2] ; AX = 200h

; Direct-Offset Addressing (doubleword array)
    mov eax, arrayD ; EAX = 10000h
    mov eax, [arrayD+4] ; EAX = 20000h
    mov eax, [arrayD+4] ; EAX = 20000h

Invoke ExitProcess, 0
main ENDP
END main
```

III. Growing Your Data (MOVZX vs MOVSX)

When you move a small piece of data into a bigger box, you have to decide what to do with the empty space.

- **MOVZX (Zero Extension):** This is the "clean" one. It just fills the empty space with zeros. It's perfect for positive numbers.
- **MOVSX (Sign Extension):** This is the "smart" one. It looks at the very first bit of your number. If it's a negative number, it fills the empty space with 1s to keep the number negative. If it's positive, it uses 0s.

IV. The Swap (XCHG)

Usually, to swap two things, you need a third hand. XCHG is like a magic trick—it swaps the values in two places at the exact same time without needing a middleman.

V. Jumping Through Arrays

Computers don't see "lists"; they see a long line of bytes.

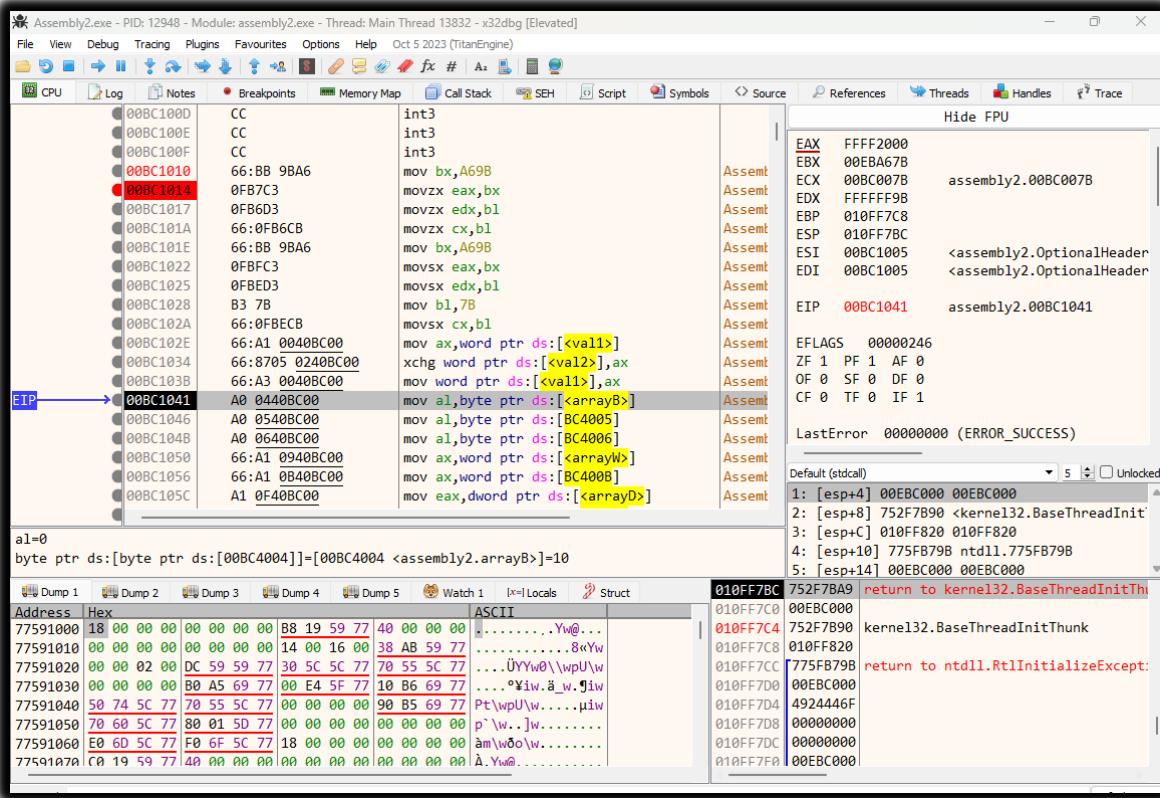
- To get to the next **Byte**, you move **1** step.
- To get to the next **Word**, you move **2** steps.
- To get to the next **Doubleword**, you move **4** steps.

VI. The Exit

ExitProcess is just the "Power Off" button. It tells Windows, "I'm done, you can take your memory back now."

DEBUGGING

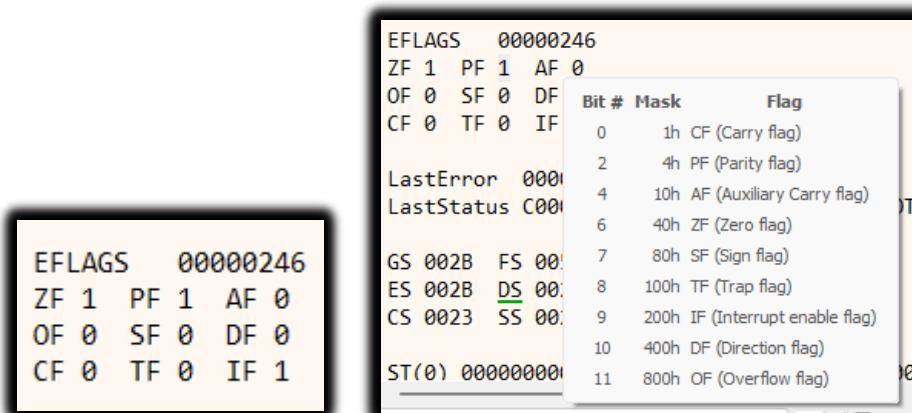
Here, we compiled the program and started debugging it in x64dbg instead of visual studio community, coz it doesn't have the registers option:



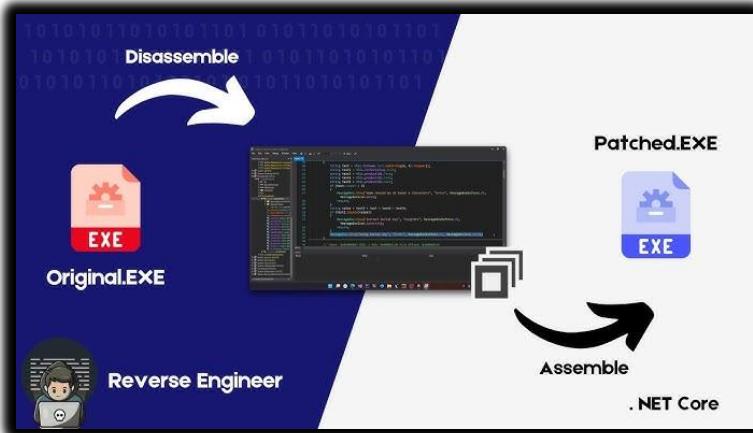
You can see the registers, flag registers, memory dumps, stack etc.

Each flag is assigned a value of 0 (clear) or 1 (set).

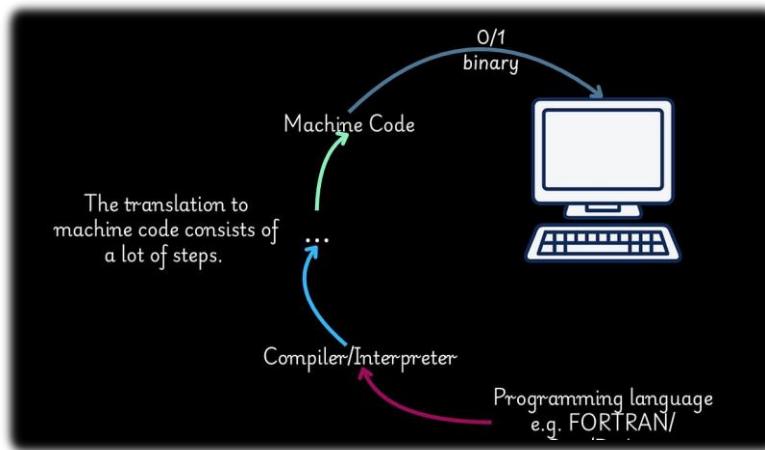
Here's an example:



Yes, it's completely normal for a disassembled EXE to have **more lines of code** than what you actually wrote.



That's because the **compiler adds extra code** to handle things like library functions, operating system calls, and other behind-the-scenes tasks.

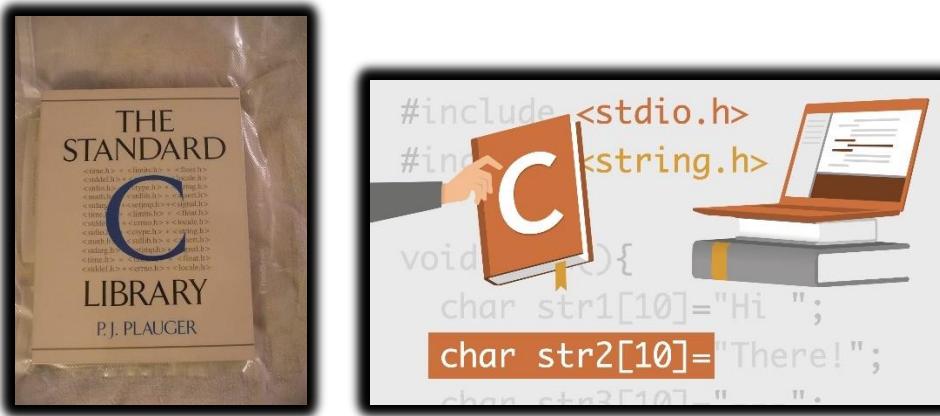


In the disassembly, you can see the compiler has added instructions to set up the stack, call `main()`, and exit the program.

It also includes code to handle errors, like division by zero or invalid memory access.



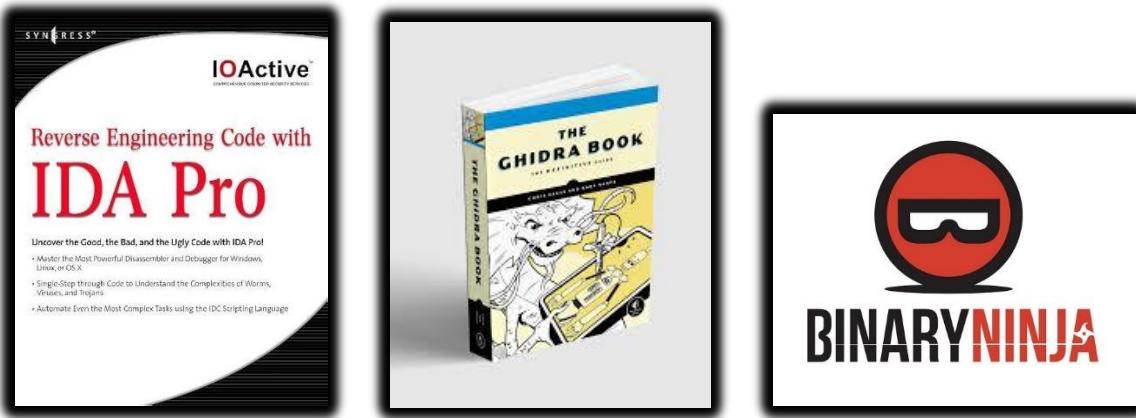
On top of that, any libraries your program uses—like the C standard library—bring in their own code, which gets included in the EXE.

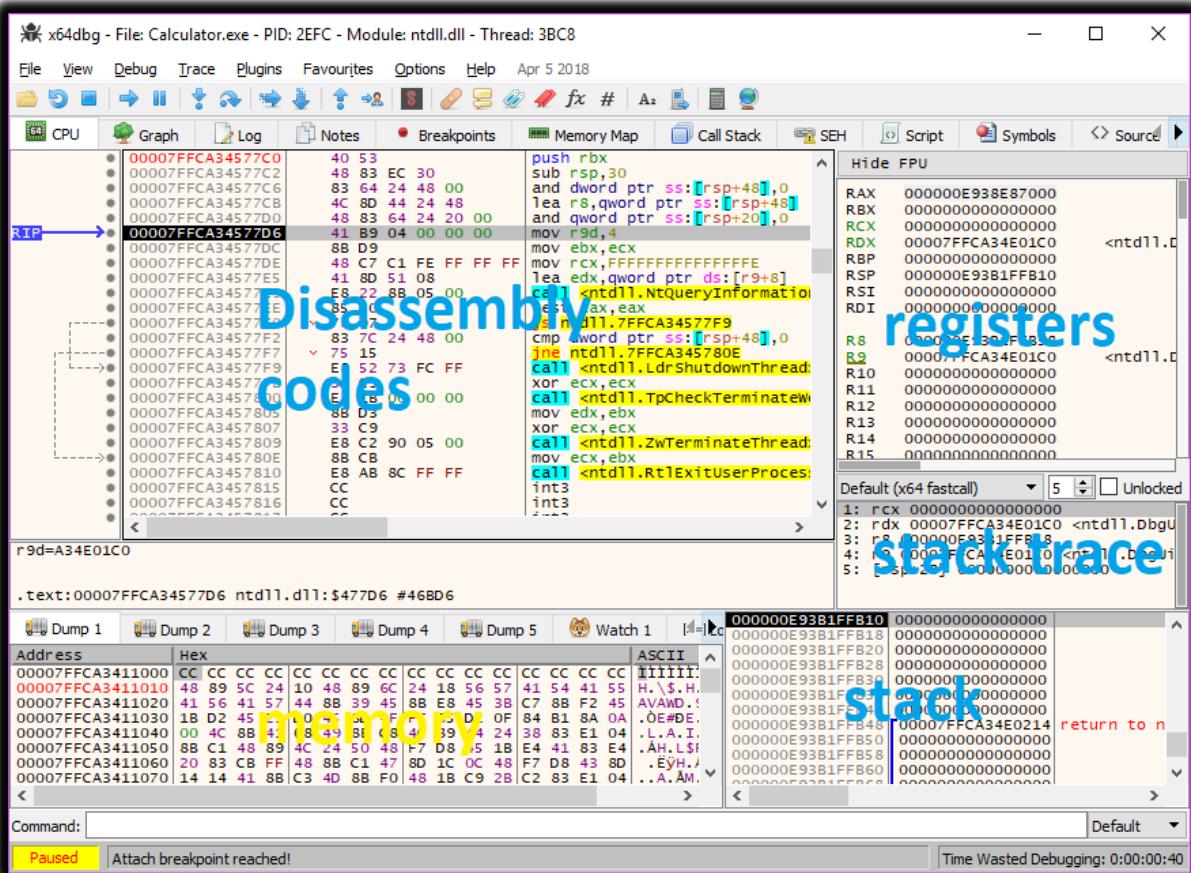


And sometimes, other programs, like debugging tools or malware, can add code too.



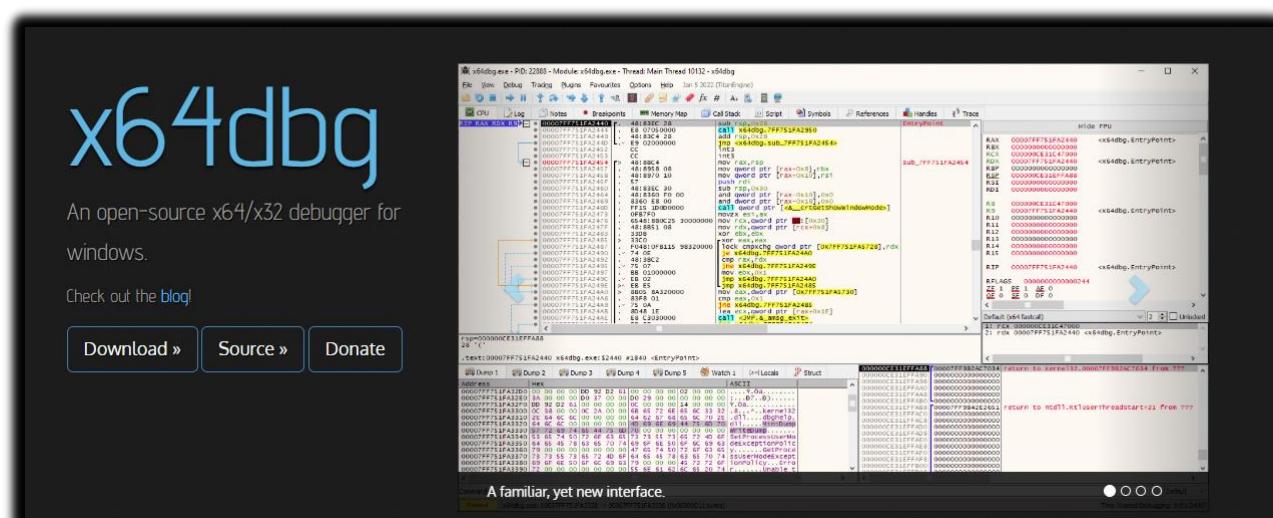
If you want to dig deeper into what's in the disassembled EXE, a debugger lets you step through the code and inspect registers and memory. You can also use a disassembler to get a detailed view of the assembly instructions.





x64dbg – Human-Friendly Window & Operand Guide

Debugging with x64dbg can feel overwhelming at first because there are a lot of windows and registers to keep track of. Let's break it down step by step, in plain English, explaining **what each window shows, why it matters, and how you can use it**.



1. Call Stack Window

What it is:

Think of the call stack as a **breadcrumb trail** of function calls that led to where the program is now.

How it works:

- Functions are listed in **reverse order**. The **top entry** is the most recent function that was called.
- You can see which function called it, and which function called that one, all the way back to the start of the program.

Why it matters:

- Helps you figure out **how you got here** in the program.
- Useful for tracking **nested function calls** and debugging crashes (like “segmentation faults”).

How to open:

View > Call Stack

2. Memory Map Window

What it is:

The Memory Map shows a **bird's-eye view of the program's memory**.

What you see:

- Loaded modules (DLLs, EXEs)
- Data segments (variables)
- Stack and heap areas

Why it matters:

- Helps you understand **where different pieces of the program are in memory**.
- Essential when working with pointers, arrays, or investigating crashes due to invalid memory access.

How to open:

View > Memory Map

3. Command Window

What it is:

This is a **text-based control panel** where you type debugger commands directly.

Things you can do:

- Start and stop the debugger
- Set and remove breakpoints
- Step through code one instruction at a time
- Inspect or change values in registers or memory

Why it matters:

- Gives you **full control** of the debugger
- Great for advanced users who want precision over what the debugger is doing

4. Registers Window

What it is:

Registers are like **tiny, ultra-fast memory slots** inside the CPU. They store:

- Current instruction pointer
- Calculation results
- Addresses for memory operations

Why it matters:

- By watching registers, you can **see exactly what the CPU is thinking**.
- Essential for understanding how the program is executing at the machine level.

5. Disassembly Window

What it is:

Shows **machine code translated into assembly instructions**.

Why it matters:

- Lets you see **what the CPU is actually doing**
- Critical for reversing programs or debugging tricky bugs

6. Variables Window

What it is:

Displays **program variables and their current values**.

Why it matters:

- Lets you track how data changes over time
- Useful for debugging logic errors or tracking unexpected program behavior

7. Threads Window

What it is:

Shows all **threads** (independent paths of execution) in a program.

Why it matters:

- Lets you **pause, resume, or switch between threads**
- Essential for debugging multithreaded programs

8. Breakpoints Window

What it is:

Lists all **breakpoints**, which are markers telling the debugger to **pause execution at specific points**.

Why it matters:

- Lets you inspect program state at **critical moments**
- Makes it easier to debug step by step instead of running blindly

Assembly Operand Basics (Human-Friendly Version)

Assembly instructions operate on **operands**, which are the “things” being manipulated. There are **three basic types**:

1. **Register operands** – stored in CPU registers (e.g., EAX, ECX)
2. **Memory operands** – stored in memory locations (e.g., [var1])
3. **Immediate operands** – fixed values built into the instruction itself (e.g., 42)

MOV Instruction Tips

1. Destination and Source:

- The **first operand** = destination (where the data goes)
- The **second operand** = source (where the data comes from)

2. Limitations:

- Destination **cannot** be a segment register or the instruction pointer (EIP)
- You can't move data **directly into special registers** like EIP

Operand Notation

- reg/mem32 → can be a **32-bit register or a 32-bit memory location**
- imm16 → a **16-bit immediate value** (a number hard-coded into the instruction)

ADDITION AND SUBTRACTION

I. INC (Increment)

The INC instruction increments the value of a register or memory operand by 1. Here's the syntax:

```
INC reg/mem
```

For example, if you have a data segment with a variable myWord initialized to 1000h, you can use INC like this:

```
.data  
myWord WORD 1000h  
  
.code  
inc myWord ; myWord = 1001h
```

II. DEC (Decrement)

The DEC instruction decrements the value of a register or memory operand by 1. Its syntax is similar to INC:

```
DEC reg/mem
```

For instance, you can use DEC to decrement the value stored in the bx register:

```
.data
myWord WORD 1000h

.code
inc myWord ; myWord = 1001h
mov bx, myWord ; Load myWord into bx
dec bx ; bx = 1000h
```

III. Flags Affected by INC and DEC

When you use the INC (increment) or DEC (decrement) instructions, they don't just change a number—they also play with the CPU's status flags. Here's what happens:

- **Overflow Flag (OF):** Changes if the operation causes a signed overflow.
- **Sign Flag (SF):** Reflects the new sign (positive or negative) of the result.
- **Zero Flag (ZF):** Turns on if the result ends up being zero.
- **Auxiliary Carry Flag (AF):** Updates for certain lower-nibble carries—useful for BCD arithmetic.
- **Parity Flag (PF):** Indicates whether the number of set bits in the result is even or odd.

Fun twist: The **Carry Flag (CF)** doesn't budge! That's a little unexpected since we usually associate arithmetic operations with the carry. So, INC and DEC can change a lot of flags, but CF stays put.

 **Pro tip:** In assembly, tiny details matter. One increment or decrement can subtly shift the CPU's state. Knowing exactly which flags are affected helps you predict behavior and write programs that behave like clockwork.

ADD INSTRUCTION

The ADD instruction is used to add a source operand to a destination operand of the same size. Here's the syntax:

```
ADD dest, source
```

or

```
add dest, source
```

Assembler doesn't care about capital letters!

- The source operand remains unchanged.
- The sum of the operands is stored in the destination operand.

For example, let's add two 32-bit integers:

```
.data
    var1 DWORD 10000h
    var2 DWORD 20000h

.code
    mov eax, var1    ;EAX = 10000h
    add eax, var2    ;EAX = 30000h
```

The instruction affects various CPU flags, including Carry, Zero, Sign, Overflow, Auxiliary Carry, and Parity. How these flags change depends on the result placed in the destination operand.

SUB INSTRUCTION

The SUB instruction subtracts a source operand from a destination operand. The syntax is the same as for ADD:

```
SUB dest, source
```

- Like ADD, the source operand remains unchanged.
- The result of the subtraction is stored in the destination operand.

For example, let's subtract two 32-bit integers:

```
.data
var1 DWORD 30000h
var2 DWORD 10000h

.code
mov eax, var1      ; EAX = 30000h
sub eax, var2      ; EAX = 20000h
```

Again, this instruction affects CPU flags such as Carry, Zero, Sign, Overflow, Auxiliary Carry, and Parity based on the value stored in the destination operand.

The SUB instruction subtracts var2 from var1. So, in this case:

- var1 is the source operand, and it contains the value 30000h.
- var2 is the destination operand, and it contains the value 10000h.

The SUB instruction subtracts var2 from var1, resulting in EAX being set to 20000h. Therefore, var2 is subtracted from var1.

NEG INSTRUCTION

The NEG (negate) instruction reverses the sign of a number by converting it to its two's complement. It can be applied to registers or memory. Here's the syntax:

To find the two's complement, reverse all the bits in the destination operand and add 1.

For example, to negate the value in EAX:

```
NEG reg/mem  
  
neg eax ; Negate the value in EAX
```

As with ADD and SUB, the NEG instruction also affects CPU flags based on the result.

IMPLEMENTING THE ARITHMETIC EXPRESSIONS

With the ADD, SUB, and NEG instructions, you can implement arithmetic expressions in assembly language.

You can break down an expression into individual operations and combine them.

For instance, if you want to calculate **Rval = -Xval + (Yval - Zval)**, you can:

```
.data  
Rval SDWORD ? ; Define a signed 32-bit variable for Rval  
Xval SDWORD 26 ; Initialize Xval to 26  
Yval SDWORD 30 ; Initialize Yval to 30  
Zval SDWORD 40 ; Initialize Zval to 40  
  
.code  
; Step 1: Negate Xval and store it in EAX  
mov eax, Xval ; EAX = 26  
neg eax ; Negate EAX, now EAX = -26  
  
; Step 2: Subtract Zval from Yval and store it in EBX  
mov ebx, Yval ; EBX = 30  
sub ebx, Zval ; Subtract Zval from EBX, now EBX = -10  
  
; Step 3: Add the results (EAX and EBX) and store it in Rval  
add eax, ebx ; Add EAX and EBX, result in EAX  
mov Rval, eax ; Store the result in Rval  
  
; At this point, Rval contains the desired value: Rval = -Xval + (Yval - Zval)
```

NEGATING A VARIABLE WITHOUT REGISTERS IN ASM

Good news: **you don't need a separate register just to negate a variable.** You can do it directly!

In my earlier example, I used extra registers to make the steps super clear—but in real-world coding, you can streamline things for efficiency.

Here's a **leaner version**: it negates Xval and calculates Rval without using any extra register:

```
.data
Rval SDWORD ? ; Define a signed 32-bit variable for Rval
Xval SDWORD 26 ; Initialize Xval to 26
Yval SDWORD 30 ; Initialize Yval to 30
Zval SDWORD 40 ; Initialize Zval to 40

.code
; Negate Xval directly
neg Xval           ; Negate Xval in place

; Calculate Rval
mov eax, Yval      ; EAX = 30
sub eax, Zval      ; Subtract Zval from EAX, now EAX = -10
add eax, Xval      ; Add Xval to EAX, now EAX = -36
mov Rval, eax       ; Store the result in Rval
```

 **Tip:** Using fewer registers keeps your code clean and can make it run faster. But sometimes, extra registers help with readability—so pick what works best for your situation.

SIGNED NUMBERS

Signed numbers let us represent **both positive and negative values** in binary. They do this by reserving **one bit for the sign** and using the rest to show the number's magnitude (absolute value). The most common system for this is **Two's Complement**.

Here's the magic behind it:

- **Sign Bit:** The leftmost bit (most significant bit) decides the sign.
 - 0 → positive
 - 1 → negative
- **Magnitude Bits:** The remaining bits show the number's size.

Negating a number in Two's Complement is simple:

1. Flip all the bits (turn 0s → 1s and 1s → 0s).
2. Add 1 to the result.

For example: Let's find out the 2's complement of given 8-bit number

$$\begin{array}{r} 00101001 \\ \text{Invert the bits} \\ 11010110 \\ + 00000001 \\ \hline 11010111 \end{array}$$

Then, add 1

The 2's complement of 00101001 is **11010111**

Examples:

- 0010 → +2 (positive, because the sign bit is 0)
- 1010 → -2 (negative, sign bit is 1)
 - Step 1: Invert all bits → 0101
 - Step 2: Add 1 → 0101 + 1 = 0110 → -2

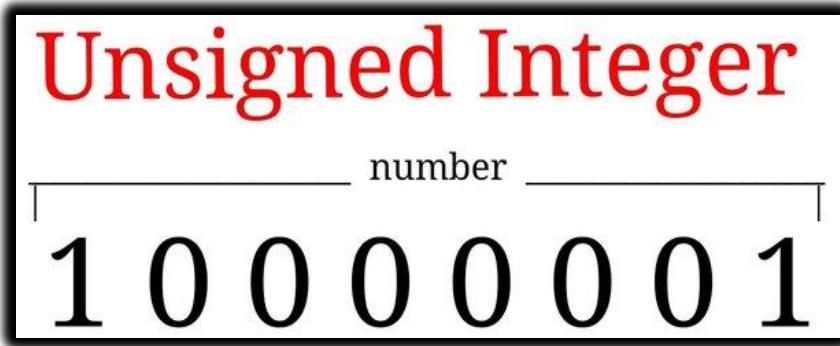
💡 Tip: Think of Two's Complement as a clever trick: it lets addition and subtraction work seamlessly, even with negative numbers, without needing separate rules.

UNSIGNED NUMBERS

Unsigned numbers are the “always-positive” kind. They **only represent zero or positive values**, because they use **all their bits just to show the magnitude**—no sign bit needed.

Examples:

- $0010 \rightarrow 2$
- There’s no way to represent negative numbers here, so things like negation just don’t exist.



Quick takeaway:

- **Signed numbers:** have a sign bit, so they can go positive **or** negative.
- **Unsigned numbers:** all positive, max value bigger for the same number of bits.

💡 Tip: Pick signed or unsigned based on what your program actually needs—don’t waste bits on a negative that will never show up!

SIGNED VS UNSIGNED NUMBERS: AVOIDING CONFUSION

Sometimes people see a binary number starting with 1 and immediately think, “Ah! It must be a negative signed number!” 😬

This is a classic mix-up if you don’t know whether the number is signed or unsigned.

Here’s how to prevent that confusion:

1. Label it clearly

- Use **descriptive variable names**: signedValue vs unsignedValue.
- Add **comments** in your code:

```
; This is unsigned: 0010 = 2  
; This is signed: 1010 = -6 (Two's Complement)
```

2. Stick to standard conventions

- Signed numbers → **Two's Complement**
- Unsigned numbers → **plain binary**
- Knowing the convention tells you how to interpret the leading bit.



3. Use the correct data types

- If your language supports it, pick int (signed) or uint (unsigned).
- Be consistent: don't mix signed and unsigned in the same calculation.



4. Document & educate

- Always **note your assumptions** about signed vs unsigned numbers.
- Teach your team: the leading 1 **doesn't automatically mean negative**—it depends on the representation.



5. Test and validate

- Run small examples to check your assumptions.
- Example: $0010 \rightarrow 2$ unsigned, $+2$ signed; $1010 \rightarrow 10$ unsigned, -6 signed.



6. Keep operations separate

- Don't mix signed and unsigned arithmetic.
- Separate them in code to avoid mistakes.



Memory Trick:

Think of signed numbers as “the leading bit is a mood indicator”:

- 0 → happy/positive

- 1 → grumpy/negative

Unsigned numbers? They’re always happy—no grumpiness allowed!

Signed value

1	1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---	---

 = -1

Unsigned value

1	1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---	---

 = 255

CPU STATUS FLAGS/ FLAG REGISTERS: WHY THEY MATTER

When your program does arithmetic, you often want to know:

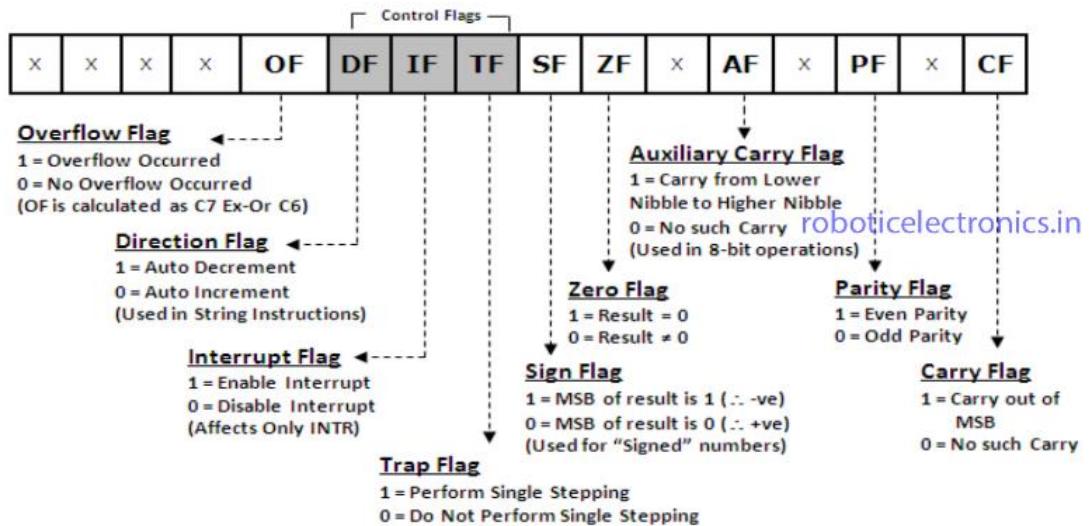
- Is the result **positive, negative, or zero?**
- Does it **fit within the range** of the variable storing it?

This is where **CPU status flags** come in—they’re tiny indicators that give you the answers instantly.

These flags are **super important**:

- They help **detect errors**.
- They let your program **make decisions** based on the result (like branching or looping).

Flag register of 8086 is a 16-bit register where status of the latest Arithmetic operation performed.



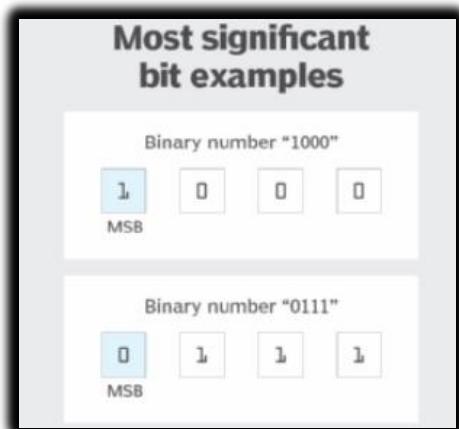
flag register of 8086

As it has 16-bits , it has 16 flags. These 16 flags are classified as

- 7 are don't care flags.
- 3 are control flags (accessible to programmers).
- 6 are status flags (not accessible to programmers).

Here's a straightforward overview of the main status flags.

First, remember this, MSB is the leftmost bit:



CARRY FLAG

Set when the result of an unsigned arithmetic operation is too large to fit into the destination operand.

$$\begin{array}{r} 1111 \text{ (Carry-In)} \\ + 1101 \\ \hline 11000 \text{ (Carry-Out)} \end{array}$$

In this example, when you add 1111 and 1101, you have a carry-out from the MSB (leftmost bit), resulting in a final carry-out value of 11000. The Carry flag would be set to indicate this carry-out condition.

For subtraction, there's typically no carry-out from the MSB, as subtraction involves taking away one value from another:

$$\begin{array}{r} 1101 \\ - 1010 \\ \hline 011 \end{array}$$

In this subtraction, there's no **carry-out from the LSB**, so the Carry flag would not be set. The carry-out occurs from the least significant bit (LSB), and propagates towards the most significant bit (MSB) during addition, not from the MSB.

This flag indicates unsigned integer overflow. For instance, if an operation has an 8-bit destination operand, but the result exceeds 255 in binary (11111111), the Carry flag is set.

→ Carry is generated

$$\begin{array}{r} 0000\ 0100 \\ 0000\ 0101 \\ \hline 0000\ 1001 \end{array}$$

Another example, subtracting 2 from 1 sets the carry flag.

FIGURE 4–4 Subtracting 2 from 1 sets the Carry flag.

0	0	0	0	0	0	0	1	(1)
+	1	1	1	1	1	1	1	0
<hr/>								(-2)
CF	1	1	1	1	1	1	1	FFh

$$CF = 1$$

$$\begin{array}{r} \overset{1}{\cancel{1}} \overset{1}{\cancel{1}} 1100 \\ + 1001 \ 0001 \\ \hline \overset{1}{\cancel{1}} \ 0100 \ 1101 \end{array}$$

$$CF = 0$$

$$\begin{array}{r} \overset{1}{\cancel{1}} \overset{1}{\cancel{1}} 1100 \\ + 0001 \ 0001 \\ \hline 1100 \ 1101 \end{array}$$

The **Carry Flag (CF)** is like the CPU's little "overflow detector" for the **most significant bit (MSB)**. It tells you whether a number got too big to fit in the bits you have.

- **CF = 1 → Carry happened**
 - This means the **sum overflowed the MSB** during an addition.
 - In plain English: the number got too big to fit, so a carry was "pushed out" of the leftmost bit.
- **CF = 0 → No carry**
 - Nothing overflowed; the result **fits perfectly** in the allocated bits.

💡 **Think of it like this:**

- CF = 1 → "Oops! Too big!" 😬
- CF = 0 → "All good, fits just right!" 😊

In short, CF helps your program **know when numbers spill over** the size limit of their binary container.

AUXILIARY CARRY FLAG

This flag is set when a 1 bit carries out of position 3 in the least significant byte of the destination operand.

Suppose we want to add 1 to the BCD value 0Fh:

```
mov al, 0Fh ; Load the value 0Fh into the AL register  
add al, 1   ; Add 1 to AL
```

Here's the binary representation and the arithmetic:

1. Binary Representation of 0Fh: 00001111
2. Binary Representation of 1: 00000001

Now, let's add them together, keeping track of the AC flag:

$$\begin{array}{r} 0000 \ 1111 \ (\text{0Fh}) \\ + 0000 \ 0001 \ (1) \\ \hline 0001 \ 0000 \ (\text{10h}) \end{array}$$

The **Auxiliary Carry (AC) flag** is a bit like the “bit 3 whisperer” of the CPU—it tells you if a carry happened **from bit 3 to bit 4** (counting from the right, starting at bit 0).

This is especially handy in **BCD (Binary-Coded Decimal) arithmetic**, where numbers are handled in a decimal-like way using binary.

Example:

- Adding 0Fh + 1 → the result is 10h.
- Notice that a **carry goes from bit 3 to bit 4**.
- That's exactly what the AC flag tracks:
 - **AC = 1** → Carry from bit 3 to bit 4 occurred
 - **AC = 0** → No carry from bit 3 to bit 4

💡 Think of it like this:

- Bit 3 is having a party. If someone “overflows” into bit 4, AC says, “Hey, bit 4 got some action!” 🎉
- If not, AC stays 0, calm and quiet.

In short, AC is a **tiny but important helper** for operations where carries between the lower nibbles matter, like in BCD arithmetic.

AF:

AF stands for auxiliary flag. As 8-bits form a byte, similarly 4 bits form a nibble. So in 16 bit operations there are 4 nibbles.

If AF = 1 ; there is a carry out from lower nibble.

If AF = 0 ;no carry out of lower nibble.

$$AF = 1$$

$$\begin{array}{r} \textcolor{blue}{1} \textcolor{blue}{1} \textcolor{red}{1} \\ 10111100 \\ + 10001001 \\ \hline 1 01000101 \end{array}$$

$$AF = 0$$

$$\begin{array}{r} \textcolor{blue}{1} \textcolor{blue}{1} \\ 10111100 \\ + 00010001 \\ \hline 11001101 \end{array}$$

PARITY FLAG (PF)

The **Parity Flag (PF)** is the CPU's way of counting 1s in the **least significant byte** (the rightmost 8 bits) of a result after an arithmetic or logical operation.

- PF = 1 → Even parity
 - The byte has an **even number of 1s**.
- PF = 0 → Odd parity
 - The byte has an **odd number of 1s**.

Think of it like this:

- PF is like a little “even-odd checker” for your bits.
- After each operation, it asks: “Do we have an even number of 1s?” 
- If yes → PF = 1 (party’s even!)
- If no → PF = 0 (party’s odd!)

It’s mostly used for **error checking** and **low-level debugging**, but knowing it can also help you understand subtle CPU behaviors.

Example 1: ADD Instruction

```
mov al, 10001100b      ; AL = 10001100  
add al, 00000010b      ; Add 00000010 to AL
```

$$\begin{array}{r} 10001100 \\ + 00000010 \\ \hline 10001110 \end{array}$$

After the ADD instruction, AL contains the binary value 10001110.

In this binary representation, there are four 0 bits and four 1 bits.

Since there is an even number of 1 bits (four), the Parity flag (PF) is set to 1.

This means that PF indicates that the result has an even parity.

Example 2: SUB Instruction

```
sub al, 10000000b      ; Subtract 10000000 from AL  
  
10001110  
- 10000000  
-----  
00001110
```

Since there is an odd number of 1 bits (one), the Parity flag (PF) is set to 0. (odd parity)

Parity Flag (PF) – Quick Example

Think of PF as the CPU's “even-odd bit detector.”

- After **ADD**, AL = 1111 (four 1s → even) → **PF = 1** ✓
- After **SUB**, AL = 0001 (one 1 → odd) → **PF = 0** ✗

💡 **In short:** PF tells you whether the **number of 1 bits** in the result is even or odd.
It's handy for **error detection** or any situation where parity matters.

PF:

It stands for parity flag.

If PF = 1 ; it means it is even parity in the result (there are even numbers of 1's).

If PF = 0 ; it means it is odd parity.

$$\text{PF} = 1$$

$$\begin{array}{r} \overset{1}{1} \ 1011\ 1100 \\ + \ 1001\ 0001 \\ \hline \overset{1}{1} \ 0100\ 1101 \end{array}$$

$$\text{PF} = 0$$

$$\begin{array}{r} \overset{1}{1} \ 1011\ 1100 \\ + \ 0001\ 0001 \\ \hline 1100\ 1101 \end{array}$$

OVERFLOW FLAG

I. Signed Numbers and Overflow

An **overflow** is set when the result of a **signed arithmetic operation** is too large or too small to fit into the destination operand.

Unsigned:

This is a 8 bit positive number which ranges from 0 to 255. In hexadecimal it's range is from 00 to FF. In the OF flag, it has nothing to do with unsigned numbers. Only signed numbers are considered in the OF flag.

Signed Numbers

A **signed number** is typically an **8-bit value** (it can also be 16-bit) that is evenly divided between **positive and negative numbers**.

1011 0011 ;is -ve

0111 1001 ;is +ve

The **Most Significant Bit (MSB)** determines the sign:

- **MSB = 0** → Positive number
- **MSB = 1** → Negative number

For example, the binary number **1011 0011** has MSB = 1, so it is negative.

Taking its **two's complement** gives **0100 1101**, which is **4D** in hexadecimal and represents **-77** in decimal.

II. Range of Signed Numbers

For an **8-bit signed number**:

- The range is **-128 to +127** (hexadecimal **80 to 7F**).
- This includes:
 - **128 positive values** (0 to +127)
 - **128 negative values** (-1 to -128)

III. Overflow Detection

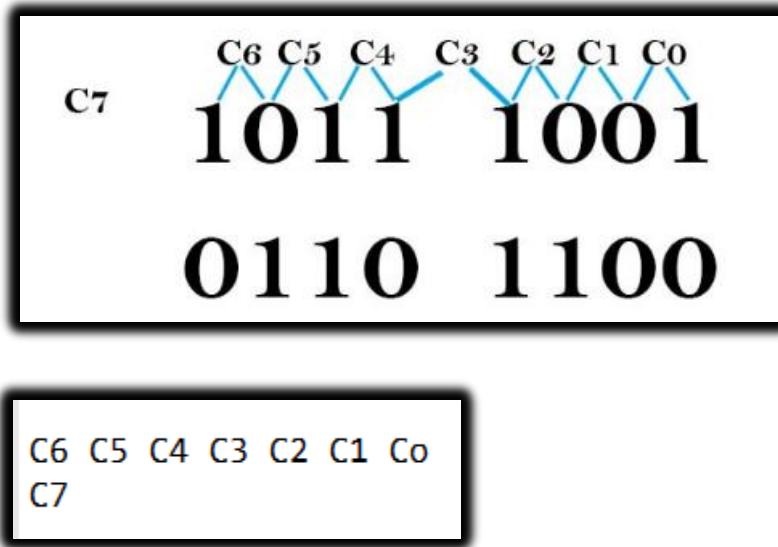
Overflow occurs when the result of an arithmetic operation **exceeds the valid signed number range**.

- If a **positive result** produces **MSB = 1**, a **positive overflow** has occurred.
- If a **negative result** produces **MSB = 0**, a **negative overflow** has occurred.

IV. Overflow Detection Mechanism

To detect overflow, the processor performs an **XOR operation** on the carry flags **C6 and C7** during arithmetic operations.

- **C0** is the carry from bit 0 to bit 1, **C1** from bit 1 to bit 2, and so on.
- If the **XOR of C6 and C7 equals 1**, it indicates that the result has exceeded the limits of signed numbers and an **overflow has occurred**.



The image represents the **carry flags from bit 0 to bit 7** that are generated during an arithmetic operation.

The **Overflow flag (OF)** is set when the result of a **signed arithmetic operation** is too large or too small to fit into the destination operand. In the case of **8-bit signed numbers**, overflow occurs when the result is:

- **Greater than +127, or**
- **Less than -128**

V. How Overflow Is Detected

To detect overflow, the processor performs an **XOR operation** on the carry flags **C6 and C7**:

- **C6** is the carry into the sign bit (bit 7)
- **C7** is the carry out of the sign bit

If the result of **C6 XOR C7 = 1**, then an **overflow has occurred**, and the Overflow flag is set.

VII. Example

Consider adding two **8-bit signed numbers**:

127 + 1

- The mathematical result is **128**
- However, 128 is **outside the valid range** of an 8-bit signed number

As a result, **overflow occurs**, and the Overflow flag is set.

Carry flag	Overflow flag
C7 = 1, C6 = 0	OF = 1

The table shows how the **carry flags (C6 and C7)** and the **Overflow flag** are affected during this addition. Since the XOR of **C6 and C7 equals 1**, the processor correctly detects an overflow condition.

VIII. Why the Overflow Flag Is Useful

The Overflow flag helps detect **errors in signed arithmetic operations**.

For example, if a program expects the result of an operation to stay within a specific signed range, it can check the Overflow flag:

- **OF = 0** → Result is valid
- **OF = 1** → Overflow occurred, and corrective action may be needed

IX. Using the Overflow Flag in Assembly Language

Here is an example of how the **Overflow flag** can be used in assembly language:

```
; Add two 8-bit numbers: 127 and 1
add al, 1

; Check the Overflow flag
jnc no_overflow

; Overflow has occurred, handle the error
overflow:
...

; No overflow, continue with the program
no_overflow:
...
```

X. Using the Overflow Flag in Assembly Language

The assembly code shown is a simple example of how the **Overflow flag (OF)** can be used to detect overflow during an addition operation. Here's how it works:

1. The instruction
add al, 1
adds the value **1** to the **AL register**.
2. After the addition, the processor updates the status flags.
The instruction **jno no_overflow** (Jump if No Overflow) transfers control to the **no_overflow** label **if the Overflow flag is not set**.
3. If the **Overflow flag is set**, execution continues to the **overflow** label, where the program can handle the error or take corrective action.

XI. Why Overflow Occurs

The Overflow flag is set when the result of a **signed arithmetic operation** is too large (or too small) to fit into the destination operand.

- In this example, the destination operand is the **AL register**, which is an **8-bit signed register**.
- The valid range for an 8-bit signed value is **-128 to +127**.
- If the result of the addition exceeds **+127** or goes below **-128**, an **overflow occurs**, and the Overflow flag is set.

XII. Key Takeaway

The **Overflow Flag (OF)** is used to detect overflow in **signed arithmetic operations**.

It plays a crucial role in ensuring correct results when working with signed data types.

The processor detects overflow by performing an **XOR operation on carry flags C6 and C7**.

If the result of this XOR operation is **1**, the Overflow flag is set, indicating that the signed result is invalid

ZERO FLAG

Set when the result of an arithmetic operation is zero.

For example, subtracting two equal values will set the Zero flag.

ZF:

This is zero flag. Whenever the output is **0** this flag is **1**.

If $ZF = 1$; output is zero.

If $ZF = 0$; output is non zero.

SIGN FLAG

The **Sign flag (SF)** is set when the result of an arithmetic operation is **negative**.

It directly reflects the value of the **Most Significant Bit (MSB)** of the destination operand.

In simple terms, the Sign flag copies the MSB of the result:

- **MSB = 1 → SF = 1**, indicating a **negative result**
- **MSB = 0 → SF = 0**, indicating a **positive result**

So, the Sign flag gives a quick way to determine whether the result of an operation is positive or negative.

I. Control of Status Flags

Status flags, including the **Sign flag**, are controlled automatically by the **Arithmetic Logic Unit (ALU)**.

This means:

- The user **cannot directly set or clear** these flags.
- Flags are updated automatically based on the result of arithmetic or logical operations.

This behavior ensures that flags always accurately represent the outcome of the last operation.

II. Using Status Flags to Detect Errors

Status flags are commonly used to **check for errors or special conditions** after arithmetic operations.

Some important examples include:

- **Overflow Flag (OF):**
If OF is set after an addition, it means the result is too large or too small to fit into the destination operand (signed overflow).
- **Zero Flag (ZF):**
If ZF is set, the result of the operation is zero. This is useful for comparisons and decision-making in programs.

III. Role of the Sign Flag in Error Detection

The **Sign flag** can also help detect unexpected results. For example:

- If two **positive numbers** are added together, the result should also be positive.
- If the **Sign flag is set** after such an operation, it indicates that something has gone wrong (such as overflow).

This makes SF useful not just for identifying negative values, but also for validating arithmetic results.

IV. Determining the Sign of a Result

The Sign flag can be directly used to determine the sign of a number:

- **SF = 1 → Result is negative**
- **SF = 0 → Result is positive**

This is especially useful in conditional branching and decision-making instructions.

V. Effect of Subtraction on the Sign Flag

The provided code demonstrates how the **Sign flag** is affected by a subtraction operation.

- First, a value is moved into a register.
- Then, a larger value is subtracted from it.
- Since the subtracted value is greater than the original value, the result becomes **negative**.
- As a result, the **Sign flag is set to 1**, indicating a negative outcome.

In the second example, subtraction is performed on an 8-bit register. The result of the subtraction is **-1**, which is represented as **FFh in hexadecimal** (two's complement form). Because the MSB of this result is 1, the **Sign flag is set**.

```
mov eax, 4
sub eax, 5 ; EAX = -1, SF = 1

mov bl, 1
sub bl, 2 ; BL = FFh (-1), SF = 1
```

VI. Final Note

The **Sign flag** is a powerful and simple tool for:

- Identifying negative results
- Validating arithmetic operations
- Supporting decision-making in programs

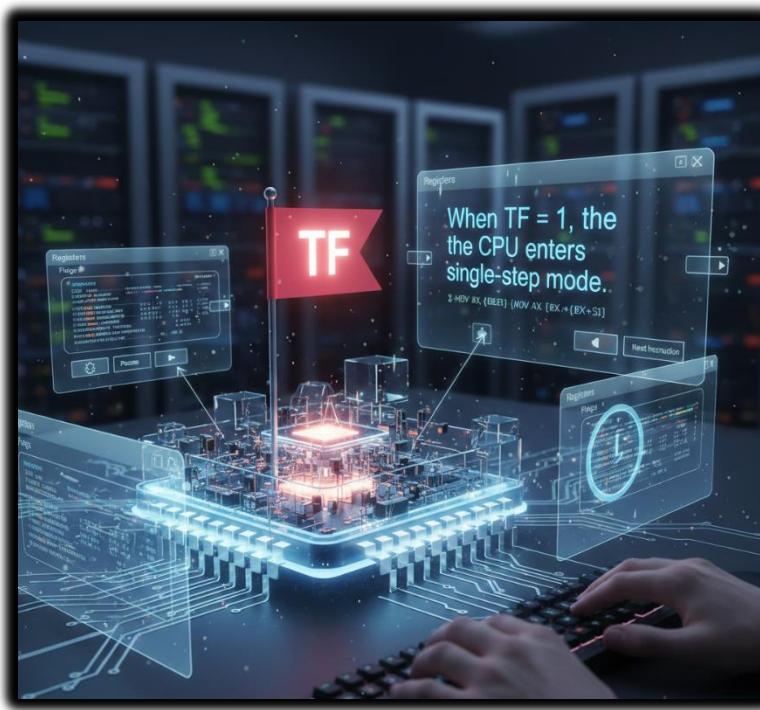
By understanding how the Sign flag works alongside other status flags, you gain a much clearer picture of how the CPU processes arithmetic operations and handles errors.

CONTROL FLAGS

Control flags are special flags that control how the CPU operates. In the **8086 microprocessor**, there are **three control flags**:

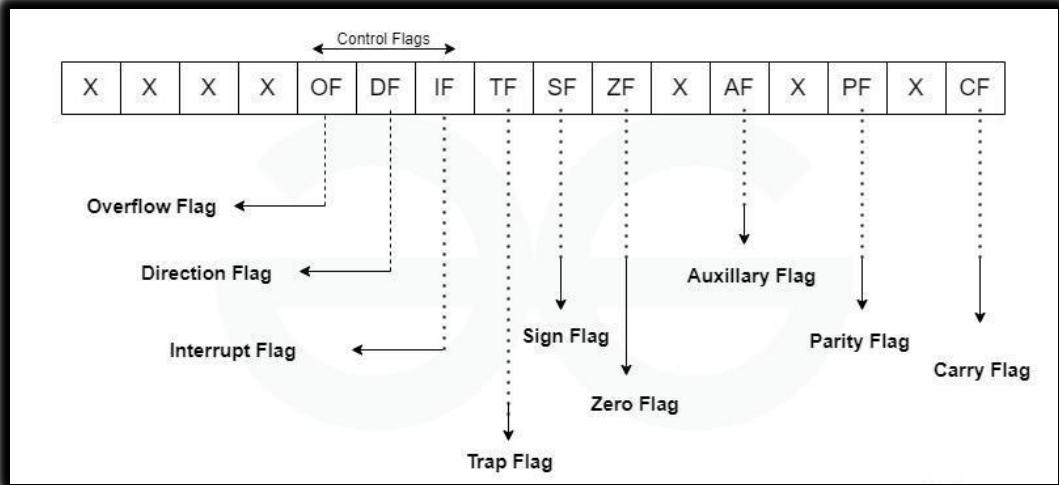
I. Trap Flag (TF)

- When **TF = 1**, the CPU enters **single-step mode**.
- In this mode, the CPU executes **one instruction at a time** and then pauses.
- This is especially useful for **debugging**, as it allows programmers to closely observe how each instruction affects registers and flags.



II. Interrupt Flag (IF)

- When **IF = 1**, the CPU **allows interrupts**.
- Interrupts are signals from **external devices** (such as keyboard, timer, or I/O devices) that require the CPU's attention.
- When an interrupt occurs, the CPU:
 1. Temporarily stops the current program
 2. Executes the interrupt service routine
 3. Returns to the original program after handling the interrupt



III. Direction Flag (DF)

- The **Direction flag** controls the direction of **string operations**.
- When **DF = 0**, the CPU **increments** the string pointer after each operation.
- When **DF = 1**, the CPU **decrements** the string pointer after each operation.

(Image showing TF, IF, and DF in the 8086 flag register)

By default, these flags are **cleared (set to 0)**, which means:

- The CPU is **not** in single-step mode
- Interrupts are **disabled**
- String operations move in the **forward direction**



IV. Uses of Control Flags

Debugging

- The **TF flag** can be used to execute a program **one instruction at a time**.
- This makes it easier to trace program flow and identify errors.
- For example, TF can help observe how string instructions behave when the **DF flag** changes.



Disabling or Enabling Interrupts

- The **IF flag** allows programmers to **disable interrupts** when needed.
- This is useful during debugging or while executing critical code sections where interruptions are undesirable.



Controlling String Operations

- The **DF flag** determines whether string operations move forward or backward in memory.
- For example, setting **DF = 1** can be useful when **reversing a string**.



V. Summary

- **TF** → Controls single-step execution (debugging)
- **IF** → Enables or disables interrupts
- **DF** → Controls the direction of string operations

Control flags play a crucial role in managing CPU behavior. Understanding how they work helps in debugging programs, handling interrupts efficiently, and performing advanced string operations.

```
;AddSubTest.asm - Arithmetic and Flags Demonstration
.386
.model flat,stdcall
.stack 4096

ExitProcess proto,dwExitCode:dword

.data
    Rval SDWORD ?
    Xval SDWORD 26
    Yval SDWORD 30
    Zval SDWORD 40

.code

main PROC
    ; INC and DEC
    mov ax, 1000h
    inc ax          ; Increment: AX = 1001h
    dec ax          ; Decrement: AX = 1000h

    ; Expression: Rval = -Xval + (Yval - Zval)
    mov eax, Xval   ; Load Xval into EAX
    neg eax         ; Negate EAX (convert to negative)
    mov ebx, Yval   ; Load Yval into EBX
    sub ebx, Zval   ; Subtract Zval from Yval
    add eax, ebx    ; Add the negated Xval and (Yval - Zval) to get Rval
    mov Rval, eax   ; Store the result in Rval

    ; Zero Flag example
    mov cx, 1
    sub cx, 1       ; ZF = 1 because the result is zero

    ; Sign Flag example
    mov cx, 0
    sub cx, 1       ; SF = 1 because the result is negative
    mov ax, 7FFFh
    add ax, 2        ; SF = 1 because the result is negative

    ; Carry Flag example
    mov al, 0FFh
    add al, 1        ; CF = 1, AL = 00 (overflow)

    ; Overflow Flag example
    mov al, 127      ; Maximum positive value
    add al, 1        ; OF = 1 (overflow)
    mov al, -128     ; Minimum negative value
    sub al, 1        ; OF = 1 (overflow)

    INVOKE ExitProcess, 0
main ENDP

END main
```

OFFSET

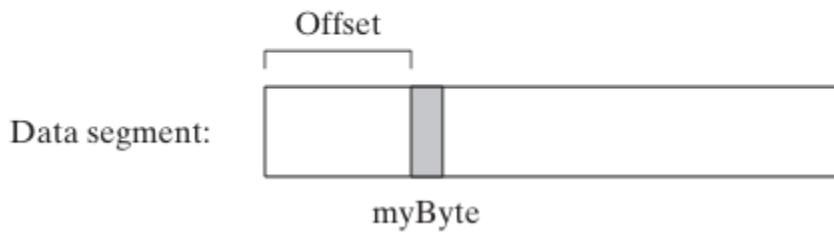


The image shows myByte variable, an offset, and the data segment.

The data segment is a section of memory that stores program data eg variables, constants, strings.

The offset is the distance from the beginning of the data segment to the variable myByte.

A variable named myByte.



I. Data Segment, Offset, and Memory Addressing

If the **data segment** starts at memory address **1000h** and the **offset** of a variable myByte is **200h**, then the variable will be stored at:

$$1000h + 200h = 1200h$$

So, myByte is located at memory address **1200h**.

II. What Is an Offset?

The **offset** is the distance of a variable from the **start of its segment**.
The processor uses this offset to locate variables in memory.

When the CPU needs to access myByte, it:

1. Takes the **starting address of the data segment**
2. Adds the **offset of myByte**
3. Uses the result as the **memory address**

The **operating system** keeps track of:

- The starting address of each program's data segment
- The size of that data segment

This helps protect one program's data from being accessed by another program.

IV. Correct Assembly Example (8086-style)

Here is a **clean and correct** example showing how to access myByte:

```
mov ax, data_segment      ; Load the starting address of the data segment
add ax, offset myByte    ; Add the offset of myByte
mov al, [ax]              ; Load the value of myByte into AL
```

V. How This Code Works (Step-by-Step)

1. **mov ax, data_segment**
Loads the starting address of the data segment into the **AX register**.
2. **add ax, offset myByte**
Adds the offset of myByte to AX, giving the **exact memory address** of the variable.
3. **mov al, [ax]**
Loads the value stored at that memory address into **AL**.

The **[ax]** notation means “use the memory location whose address is stored in AX”.
This is called **dereferencing** the register.

After execution, **AL contains the value of myByte**.

Is the Statement About Offset Correct?

Yes — it is correct. 

The offset is added to the data segment's starting address to get the memory address of a variable. This is exactly how the **8086 microprocessor** accesses data in memory.

VI. Key Points About Offsets

- The offset is a **16-bit value**
- It allows access to up to **64 KB of memory**
- This matches the **maximum size of a data segment** in the 8086

VII. Instructions That Use Offsets

- **mov** — moves data between registers and memory
- **lea** — loads the effective address (offset) of a variable into a register

Example:

```
lea ax, myByte ; Loads the offset of myByte into AX
```

VIII. Final Summary

- A variable's **memory address = data segment address + offset**
- The offset tells the CPU *where the variable is located inside the segment*
- [ax] means "*read the value from the memory address stored in AX*"
- Understanding offsets is essential for working with **segments, pointers, and memory** in assembly language

HOW OFFSETS WORK IN ASSEMBLY

```
.data  
    bigArray DWORD 500 DUP(?)  
    pArray    DWORD bigArray  
  
    mov esi, pArray  
    ;At this point, ESI contains the address of the beginning of the bigArray
```

The assembly statement: pArray DWORD bigArray, defines a **doubleword (DWORD) variable** named **pArray** and initializes it with the **address of bigArray**.

In simple terms, pArray acts as a **pointer** that points to the **first element** of the bigArray.

The array **bigArray** contains **500 doubleword (DWORD) elements**, so pArray holds the starting address of this block of memory.

I. Using the OFFSET Operator

The **OFFSET operator** is commonly used to obtain the **offset address** of a variable in memory.

When accessing a variable, the processor adds this offset to the **starting address of the data segment** to calculate the variable's actual memory address.

In other words:

Memory Address = Data Segment Base + OFFSET of the Variable

II. Accessing a Variable Using OFFSET

To access a variable like bVal in memory, you first load its offset and then read the value from that address.

```
.data  
    bVal BYTE ?  
    mov esi, OFFSET bVal  
    ; At this point, ESI contains the address of the variable bVal.  
    mov al, [esi]  
    ; At this point, AL contains the value of the variable bVal.
```

When the code executes:

- The OFFSET operator provides the location of bVal within the data segment.
- The processor uses this offset, along with the data segment's base address, to locate bVal in memory.
- The value stored at that memory location can then be accessed or moved into a register.

III. Dereferencing and the ESI Register

Dereferencing means **accessing the data stored at a memory address**.

In this case, the address is stored in the **ESI register**, and the CPU uses that address to read the actual data from memory.

```
.data
    bVal BYTE ?
    wVal WORD ?
    dVal DWORD ?
    dVal2 DWORD ?

.code
    mov esi, OFFSET bVal ; ESI = 00404000h
    mov esi, OFFSET wVal ; ESI = 00404001h
    mov esi, OFFSET dVal ; ESI = 00404003h
    mov esi, OFFSET dVal2 ; ESI = 00404007h
```

(If this feels confusing, revisiting pointer book in my C Github repo can really help.)

IV. Effect of the MOV Instruction with OFFSET

After the code executes, the **ESI register contains the offset** of the variable specified in the MOV instruction.

For example:

- After the **first MOV instruction**, ESI contains the offset of the variable **bVal**. This offset is **00404000h**, which is the memory location where bVal is stored.
- After the **second MOV instruction**, ESI contains the offset of the variable **wVal**. This offset is **00404001h**, which is the memory location where wVal is stored.
- The same pattern continues for subsequent variables:
each MOV instruction loads the **offset of the specified variable** into ESI.

Once the offset is in ESI, the processor can **dereference** it (for example, using [ESI]) to access the actual data stored at that address.

V. Why the OFFSET Operator Is Important

The **OFFSET operator** is a powerful and commonly used tool in assembly language. It is especially useful for:

- Accessing variables in memory
- Initializing pointer registers like ESI
- Working with arrays, structures, and buffers

In short, OFFSET gives you the **address of a variable**, and dereferencing that address lets you access the **value stored there**.

VI. Accessing an Array Element Using OFFSET and ESI

```
.data  
    myArray WORD 1,2,3,4,5  
  
.code  
    mov esi, OFFSET myArray + 4
```

The first line of code in the **data segment** declares an array named **myArray** that contains **five 16-bit (WORD) values**.

In the **code segment**, the next instruction moves the **offset of myArray plus 4** into the **ESI register**.

The **OFFSET operator** returns the offset of a variable or label, which is the distance from the beginning of the data segment to that variable or label.

In this case, since myArray is the first variable in the data segment, its offset is **0000h**.

Each element in myArray is a **WORD (2 bytes)**:

- First element → offset **0000h**
- Second element → offset **0002h**
- Third element → offset **0004h**

By adding **4** to the offset of myArray, we obtain the offset of the **third integer** in the array.

VII. Using ESI as a Pointer

The **ESI register** is a general-purpose register that can store addresses.

By loading the offset of the third element of myArray into ESI, we are effectively creating a **pointer to the third integer** in the array.

Once ESI holds this address, it can be used to access the value stored at that location.

For example, the following instruction loads the value of the third integer in myArray into the **AL register**:

```
mov al, [esi]
```

VIII. Calculating offsets

For arrays...

Array Type	Array Element Size	Offset
Byte array	1 byte	+1
Word array	2 bytes	+2
Doubleword array	4 bytes	+4
Quadword array	8 bytes	+8

BYTE Arrays:

If you are working with a BYTE array, where each element is 1 byte in size, you can access the next element by adding +1 to the current element's address.

```
; Assuming byteArray starts at address 00400000h
mov esi, OFFSET byteArray ; Load the starting address of byteArray
add esi, 1 ; Move to the second element (1 byte offset)
mov al, [esi] ; Dereference/Load the value of the second element into AL
```

WORD Arrays:

For WORD arrays, where each element is 2 bytes (16 bits) in size, you should add +2 to move to the next element's address.

```
; Assuming wordArray starts at address 00400000h
mov esi, OFFSET wordArray ; Load the starting address of wordArray
add esi, 2 ; Move to the third element (2 bytes offset)
mov ax, [esi] ; Load the value of the third element into AX
```

DWORD Arrays:

DWORD arrays have elements that are 4 bytes (32 bits) in size. To access the next element, add +4 to the current element's address.

```
; Assuming dwordArray starts at address 00400000h
mov esi, OFFSET dwordArray ; Load the starting address of dwordArray
add esi, 4 ; Move to the fourth element (4 bytes offset)
mov eax, [esi] ; Load the value of the fourth element into EAX
```

QWORD Arrays:

If you are dealing with QWORD arrays, where each element is 8 bytes (64 bits) in size, you should add +8 to move to the address of the next element.

```
; Assuming qwordArray starts at address 00400000h
mov esi, OFFSET qwordArray ; Load the starting address of qwordArray
add esi, 8 ; Move to the fifth element (8 bytes offset)
mov rax, [esi] ; Load the value of the fifth element into RAX
```

ALIGN DIRECTIVE

The ALIGN directive is used to place a variable or label on a specific memory boundary, such as a byte, word, doubleword, or paragraph boundary.

Syntax: **ALIGN bound**

Where bound can be **1, 2, 4, or 16**.

For example:

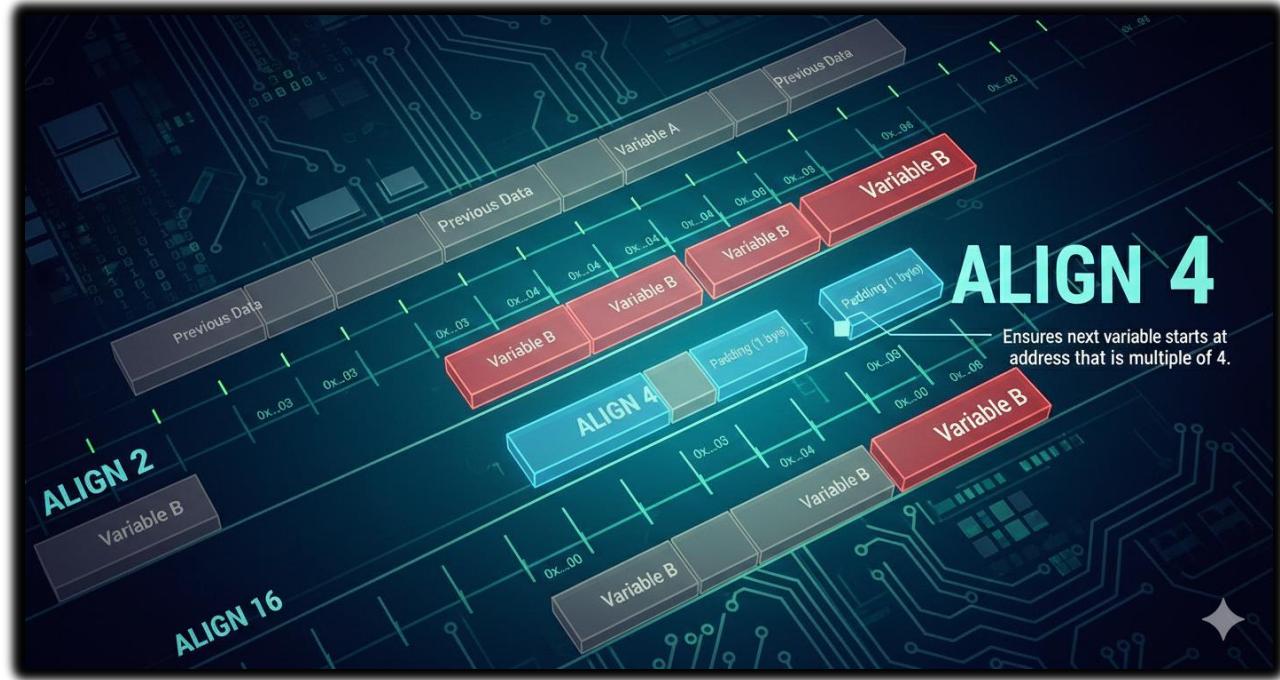
- ALIGN 1 → byte boundary
- ALIGN 2 → word boundary
- ALIGN 4 → doubleword boundary
- ALIGN 16 → paragraph boundary

Using ALIGN 4 ensures that the next variable starts at a memory address that is a **multiple of 4**.

I. Why ALIGN Is Important

Proper alignment can **improve performance**, especially for certain instructions and hardware units.

For example, the **FPU (Floating Point Unit)** works more efficiently when data is aligned on a **doubleword boundary**.



II. PTR Operator and Data Types

The **PTR operator** is used when moving data between **memory and registers** to explicitly tell the assembler **how much data** should be accessed.

Common data types used with PTR include:

- BYTE, SBYTE
- WORD, SWORD
- DWORD, SDWORD
- FWORD, QWORD, TBYTE

The PTR operator ensures that the processor interprets the data **with the correct size**, especially when the variable's declared type does not match the destination register size.

III. Moving Smaller Data into Larger Registers

Sometimes, you may want to move **smaller memory values** into a **larger register**. This is where the PTR operator becomes very useful.

Example

```
.data  
wordList WORD 5678h, 1234h  
  
.code  
mov eax, DWORD PTR wordList ; EAX = 12345678h
```

What's Happening Here?

- wordList contains **two WORD values**
- DWORD PTR tells the assembler to treat the first **two WORDs (4 bytes)** as a **single doubleword**
- These 4 bytes are loaded into **EAX**, combining them into one 32-bit value

IV. ALIGN vs PTR — Not the Same Thing

Even though they're often mentioned together, **ALIGN and PTR serve completely different purposes**:

ALIGN (Memory Placement)	PTR (Data Access/Sizing)
<ul style="list-style-type: none">• Controls where data is placed in memory.	<ul style="list-style-type: none">• Controls how much data is accessed.
<ul style="list-style-type: none">• Works at assembly time (directive).	<ul style="list-style-type: none">• Affects data interpretation (override).
<ul style="list-style-type: none">• Improves performance and alignment.	<ul style="list-style-type: none">• Prevents size-mismatch errors.

Example of ALIGN Directive:

```
.data  
myData BYTE 10          ; Byte-sized variable  
  
ALIGN 4                ; Align next variable on a 4-byte boundary  
myDouble DWORD 29       ; Doubleword-sized variable
```

Here:

- ALIGN 4 ensures myDouble starts at an address divisible by 4
- This improves access efficiency and guarantees proper alignment

V. Final Takeaway

- **ALIGN** decides *where* data lives in memory
- **PTR** decides *how big* the data is when you access it
- They solve **different problems** but are both essential for clean, efficient assembly code

BYTE PTR Operator

Think of **BYTE PTR** as a way to tell the computer, "I know this variable is big, but I only want to look at one byte of it right now."

By default, if you have a large variable (like a 4-byte Doubleword), the assembler expects you to move all 4 bytes at once.

Using the **PTR** operator lets you override that. It's like using a magnifying glass to pick out a specific "chunk" of data rather than grabbing the whole box.

You'll see this a lot when moving data between memory and registers—it's basically a size-check to make sure the data fits perfectly into the register you're using.

Here is how you'd use **BYTE PTR** to pluck a single byte out of a larger **myDouble** variable:
myDouble is a **doubleword(4 bytes)**, so the first byte is 78h.

```
.data  
myDouble DWORD 12345678h  
  
.code  
mov bl, BYTE PTR myDouble ; BL = 78h
```

This code will move the value of the least significant byte of the variable myDouble into the register BL.

Memory Layout of the myDouble Variable. The following image shows the memory layout of the myDouble variable:

Doubleword	Word	Byte	Offset	
12345678	5678	78	0000	myDouble
		56	0001	myDouble + 1
	1234	34	0002	myDouble + 2
		12	0003	myDouble + 3

The least significant byte of the variable myDouble is at offset 0000h. The most significant byte of the variable myDouble is at offset 0003h.

I. Moving Smaller Values into Larger Destinations

We can use the **BYTE PTR** operator to move smaller values into larger destination operands.

For example, the following code moves two words from memory to the register EAX:

```
.data  
wordList WORD 5678h, 1234h  
  
.code  
mov eax, DWORD PTR wordList
```

II. Grabbing Data with PTR

This code will move the two words from the wordList array into the lower half and upper half of the register EAX, respectively.

```
.data  
    myValue DWORD 12345678h ; Declare a doubleword-sized variable  
  
.code  
    mov eax, DWORD PTR myValue ; Load the entire DWORD into EAX  
    mov al, BYTE PTR myValue   ; Load only the lowest byte into AL
```

Think of PTR as a size filter.

If you have a big 32-bit variable (a **DWORD**), the computer usually wants to grab the whole thing. PTR lets you change the rules:

- **DWORD PTR**: You're saying, "Give me the whole 4-byte chunk." This fills the entire EAX register.
- **BYTE PTR**: You're saying, "Just give me the very first byte (the 'LSB')." This fits perfectly into the tiny AL register.

III. ALIGN vs. PTR: What's the difference?

They aren't even doing the same job.

1. **ALIGN is for Organization**: Imagine you're parking cars. ALIGN makes sure every car starts exactly at a line so they're easy to find and the computer can "drive" to them faster. It lives in your data section.
2. **PTR is for Execution**: This is used inside your code (the instructions). It tells the computer exactly how much data to "pick up" right now—whether it's a tiny 1-byte crumb or the whole 4-byte cookie.

TLDR: ALIGN organizes the shelf; PTR decides how much you're taking off the shelf.

TYPE, LENGTHOF, SIZEOF, LABEL

The TYPE operator returns the size, in bytes, of a single element of a variable.

Think of the **TYPE** operator as a quick way to ask the computer, "How much space does one single item in this list take up?"

Instead of you having to remember if a variable is a **Byte** or a **DWORD**, you just use TYPE. It returns a simple number representing the byte count:

- **Byte:** Returns **1**
- **Word:** Returns **2**
- **Doubleword (DWORD):** Returns **4**
- **Quadword (QWORD):** Returns **8**

I. Why use it?

It's much safer than hard-coding numbers. If you ever change your array from WORD to DWORD, using TYPE means your code updates itself automatically without you having to hunt down every "+2" and change it to a "+4."

II. The Rest of the "Size" Family:

- **LENGTHOF:** Counts how many **items** are in the array.
- **SIZEOF:** Multiplies LENGTHOF by TYPE to give you the **total bytes** the whole array uses.
- **LABEL:** Gives a name to a memory location without actually taking up any space itself.

```
.data
var1 BYTE ?
var2 WORD ?
var3 DWORD ?
var4 QWORD ?

;The following table shows the value of each TYPE expression.

Expression  Value
TYPE var1    1
TYPE var2    2
TYPE var3    4
TYPE var4    8
```

Think of **LENGTHOF** as a simple counter that tells you how many "items" are in a list, rather than the total number of bytes.

- **It's smart with DUP:** If you tell the computer to repeat an item 50 times using 50 DUP(0), **LENGTHOF** knows there are 50 items there.
- **The "First Line" Catch:** This is the part that trips people up. If you start an array on one line and continue it on the next line (without using a comma or specific formatting), **LENGTHOF** only counts what it sees on that very first line. It's like it stops reading the list as soon as it hits the end of the row.

III. The Cheat Sheet

If you have myArray WORD 10, 20, 30:

- **TYPE** is 2 (because it's a Word).
- **LENGTHOF** is 3 (because there are 3 numbers).
- **SIZEOF** is 6 (3 items \times 2 bytes each).

```
.data
byte1 BYTE 10,20,30
array1 WORD 30 DUP(?)
array2 WORD 5 DUP(3 DUP(?))
array3 DWORD 1,2,3,4
digitStr BYTE "12345678",0

; The following table lists the values returned by each LENGTHOF expression:
Expression      Value
LENGTHOF byte1    3
LENGTHOF array1   30
LENGTHOF array2   5 * (3 * 3)
LENGTHOF array3   4
LENGTHOF digitStr 9
```

```
.data
byte1 BYTE 10,20,30
array1 WORD 30 DUP(?)
array2 WORD 5 DUP(3 DUP(?))

.code
mov ax, LENGTHOF byte1 ; AX = 3
mov ax, LENGTHOF array1 ; AX = 30
mov ax, LENGTHOF array2 ; AX = 15
```

IV. SIZEOF OPERATOR

The SIZEOF operator returns a value that is equivalent to multiplying LENGTHOF by TYPE.
For example:

```
.data  
    intArray WORD 32 DUP(0)  
  
.code  
    mov eax, SIZEOF intArray ; EAX = 32 * 2 = 64
```

IntArray is indeed an array of 32 words. Each word is 16 bits or **2 bytes** in size. Therefore, the correct calculation for SIZEOF intArray should be:

$$\text{SIZEOF intArray} = 32 * 2 = 64$$

So, the correct value for EAX should be 64.

I know for beginners your brains have bluescreened right here, let me re-explain properly.

SIZEOF, LENGTHOF, and TYPE — Explained Like a Human

At first glance, SIZEOF, LENGTHOF, and TYPE sound intimidating... but they're really just answering **three very simple questions** about your data.

Let's break them apart **before** we put them together.

I. LENGTHOF — “How many items are there?”

LENGTHOF tells you **how many elements** are in the array.

If you have:

```
intArray WORD 32 DUP(?)
```

That means:

- intArray has **32 elements**
- Each element is one WORD

So: **LENGTHOF intArray = 32**

👉 Important:

LENGTHOF does **NOT** care about size in bytes.
It only counts **how many items** exist.

II. TYPE — “How big is ONE item?”

TYPE tells you the **size of a single element**, in bytes.

- BYTE = 1 byte
- WORD = 2 bytes
- DWORD = 4 bytes
- QWORD = 8 bytes

Since intArray is made of **WORDS**:

```
TYPE intArray = 2 bytes
```

👉 Think of TYPE as the **size of one slot**, not the whole array.

III. SIZEOF — “How much space does this take TOTAL?”

This is where beginners usually lose it 😅

SIZEOF is **NOT magic**. It literally just does this:

```
SIZEOF = LENGTHOF × TYPE
```

So, for intArray:

```
LENGTHOF = 32
TYPE      = 2 bytes
-----
SIZEOF   = 64 bytes
```

IV. Why EAX Ends Up Being 64

If your code does something like:

```
mov eax, SIZEOF intArray
```

What's happening mentally is:

"Hey assembler, how many BYTES does this entire array occupy?"

And the assembler answers: **"32 items × 2 bytes each = 64 bytes"**

So: EAX = 64

- ✓ Correct
- ✓ Logical
- ✓ No black magic involved

V. The Mental Model That FINALLY Makes It Click

Think of an egg carton \bigcirc :

- LENGTHOF → how many eggs
- TYPE → size of one egg
- SIZEOF → how much space the whole carton takes

You don't argue with the carton.

You just multiply.

VI. Why Beginners Get Wrecked Here

Because they:

- Confuse **elements** with **bytes**
- Think SIZEOF "counts items"
- Forget that WORD ≠ 1 byte

Once they realize: **"SIZEOF is always bytes"**

Everything snaps into place.

VII. One-Line Rule to Remember Forever

SIZEOF always tells you how many BYTES the data occupies.

LENGTHOF counts items.

TYPE tells you the size of one item.

That's it. No more suffering, my beginners...😊

NB: AI is a mega-brain use it as a partner, not a slop generation machine....

LABEL DIRECTIVE

The **LABEL directive** lets you give a name to a memory location and specify its size **without actually reserving any storage**.

You'll often see it used to create an alternative name for a variable in the data segment, or to attach a size attribute to a variable that's declared immediately after it.

```
.data
    val16 LABEL WORD
    val32 DWORD 12345678h

.code
    mov ax, val16 ; AX = 5678h
    mov dx, [val16+2] ; DX = 1234h
```

In this example, val16 is just another name for the same memory location as val32. The **LABEL directive** doesn't create any new storage on its own—it simply gives an alternate name and size to an existing variable.

```
.data
    val16 LABEL WORD
    val32 DWORD 12345678h

    mov ax, val16 ; AX = 5678h
    mov dx, [val16+2] ; DX = 1234h

    LongValue LABEL DWORD
    val1 WORD 5678h
    val2 WORD 1234h

    mov eax, LongValue ; EAX = 12345678h
```

The **LABEL** directive is like giving an existing variable a nickname.

You can also attach a size to that nickname without reserving any extra memory.

This is handy when you want to refer to the same data in a slightly different way—maybe because you’re working with different sizes or just want a clearer name in your code.



For example, you might have a 32-bit variable called val32, but sometimes you only need to look at it as a 16-bit value.

With LABEL, you can create a “shortcut” called val16 that points to the same spot in memory. No extra space is taken—val16 is just another name for val32.

```
.data
    val16 LABEL WORD
    val32 DWORD 12345678h

.code
    mov ax, val16 ; AX = 5678h
```

In this example, the label val16 lets you peek at just the **first two bytes** of the 32-bit variable val32.

It’s like looking at a slice of a bigger cake without cutting the whole thing.

You can also use a LABEL to **build bigger numbers from smaller ones**.

For instance, imagine you have two 16-bit variables, val1 and val2, and you want to treat them together as a single 32-bit value.

By creating a label called LongValue, you can access both of them as if they were one larger integer—again, without allocating any extra memory.

```
.data
    LongValue LABEL DWORD
    val1 WORD 5678h
    val2 WORD 1234h

.code
    mov eax, LongValue ; EAX = 12345678h
```

In this example, the LongValue label is used to access the 32-bit integer that is constructed from the val1 and val2 variables.

The LABEL directive does not actually allocate any storage.

It simply creates a label with a specific size attribute.

This can be useful for creating aliases for variables, constructing larger integers from smaller integers, and other purposes.

I. Understanding Labels and Memory Access: The LongValue Example

The **LongValue** label in assembly doesn't actually store the values of val1 and val2. Instead, it acts as a **pointer** to a memory location. Here's how it works step by step:

1. Memory Reservation

In the .data section, LongValue is declared as a LABEL DWORD. This doesn't assign a value—it just reserves **4 bytes of memory** (a doubleword) and gives it a name.

2. Declaring Values

val1 and val2 are declared as 16-bit words with values 5678h and 1234h. They are stored in memory consecutively, right after LongValue.

3. Loading Data

In the .code section, the instruction:

```
mov eax, LongValue
```

...moves **4 bytes of data** from the memory location pointed to by LongValue into the EAX register.

Since LongValue itself doesn't contain a specific value, the CPU reads the **next 4 bytes in memory**, which happen to be the bytes of val1 and val2.

II. Memory Layout Visualization

Let's assume little-endian storage (least significant byte first):

Memory Address	Data (Hex)	ASM Interpretation
[LongValue]	??	Uninitialized (?)
[LongValue + 1]	78h	
[LongValue + 2]	56h	val1: 5678h (WORD / 16-bit)
[LongValue + 3]	34h	
[LongValue + 4]	12h	val2: 1234h (WORD / 16-bit)

When you do: **mov eax, LongValue**

EAX gets loaded as: **EAX = 12345678h**

So, effectively, the consecutive bytes from memory (val1 and val2) are loaded into EAX.

The assembler doesn't need LongValue to hold actual data—it just provides a reference point in memory.

III. Key Takeaways

- Labels are memory references, not variables that must hold values.
- MOV with a label loads data from the memory starting at that address.
- In little-endian systems, lower bytes come first, so multi-byte values like words and doublewords are concatenated accordingly.

IV. Section Review: Operators Recap

Question	Answer	Explanation / Reality Check
The OFFSET operator always returns a 16-bit value.	FALSE	OFFSET returns an address . In Protected Mode (x86), it's 32-bit; in Real Mode, it's 16-bit.
The PTR operator returns the 32-bit address of a variable.	TRUE	PTR specifies the size of the operand at that address (e.g., DWORD PTR).
TYPE operator returns 4 for doubleword operands.	TRUE	TYPE returns the size in bytes of a single element (DWORD = 4).
LENGTHOF operator returns the number of bytes in an operand.	FALSE	LENGTHOF counts the number of elements (e.g., items in an array), not the byte count.
SIZEOF operator returns the number of bytes in an operand.	TRUE	The formula is: SIZEOF = LENGTHOF × TYPE .

DIRECT ADDRESSING MODE

Direct addressing is a memory addressing mode in which the **operand of an instruction is the actual address of the data** to be accessed.

This allows the CPU to reach the data **directly**, without performing any extra calculations.

Because of this, direct addressing is one of the **simplest and most efficient** addressing modes.

Limitations of Direct Addressing

While direct addressing is fast and straightforward, it does have some limitations:

I. Address Must Be Known at Compile Time

The memory address of the data must be known when the program is assembled. As a result, direct addressing **cannot be used for dynamically allocated data**, such as data read from a file or created at runtime.

II. Limited to the CPU's Address Space

Direct addressing can only access data that resides **within the CPU's addressable memory space**.

It cannot be used to directly access data stored in **external storage** (such as a hard disk).

Where Direct Addressing Is Commonly Used

Despite these limitations, direct addressing is widely used because of its simplicity and speed.

It is commonly used to access:

- Global variables
- Static data
- Constants stored in memory

Examples of Direct Addressing

Here are some examples of **direct addressing** in assembly language:

```
;Load the content of memory address 100 to register R1.  
LOAD R1, 100  
  
;Load the content of register R2 to register R1.  
LOAD R1, R2  
  
;Store the content of register R1 to memory address 200.  
STORE R1, 200
```

Direct addressing can also be used in high-level programming languages.

The following C code uses direct addressing to access an element of an array:

```
int main() {
    int array[10] = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10};

    ;Load the element at index 5 of the array to the register EAX.
    EAX = array[5];

    ;Return the value in register EAX.
    return EAX;
}
```

Think of **Direct Addressing** as telling the computer exactly where to go using a specific "home address."

When you load something like the 5th item in an array into a register, the computer doesn't have to do any math or "think" about where it is—it just goes straight to that spot in one shot. It's fast because there are no extra steps to calculate the address.

The Catch

The big downside is that you have a very limited amount of "room" to write out that address.

Imagine trying to write a really long, specific address on a tiny sticky note; eventually, you run out of space.

In the same way, direct addressing often can't handle very large or complex memory locations because the instruction itself only has so many bits to hold the address.

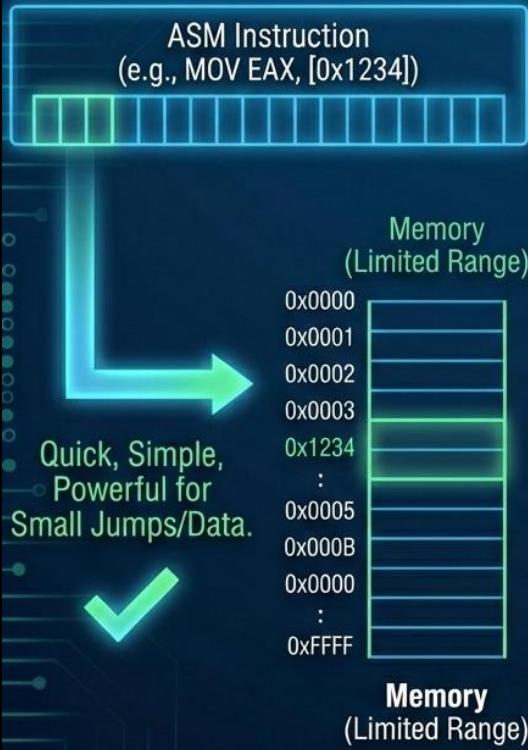
It's powerful for quick, simple jumps, but you hit a wall once your program gets too big or complex.

Quick Summary

- Direct addressing uses the **actual memory address** of the operand
- It requires **no extra calculations**
- It is fast, simple, and efficient
- Best suited for **static and known-at-compile-time data**

Direct Addressing Limitations

Small Program (Direct Addressing Works)



Large Program (Direct Addressing Hits a Wall)



Conclusion: Direct Addressing is fast but limited by the number of bits in the instruction itself. It fails once memory locations become too large or complex for the instruction's address field.

INDIRECT ADDRESSING MODE

Indirect addressing is a memory addressing mode in which the operand of an instruction is **not the data itself**, but a **pointer to the data**.

The actual address of the data is stored in a **register or memory location**, so the CPU must first **dereference** the pointer to access the data.

Because of this extra step, indirect addressing is **a bit more complex than direct addressing**, but it comes with several powerful advantages.

Advantages of Indirect Addressing

I. Dynamic Data Access

The address of the data **doesn't need to be known at compile time**.

This means indirect addressing can be used to access **data that is created at runtime**, such as:

- Data read from files
- Data dynamically allocated in memory

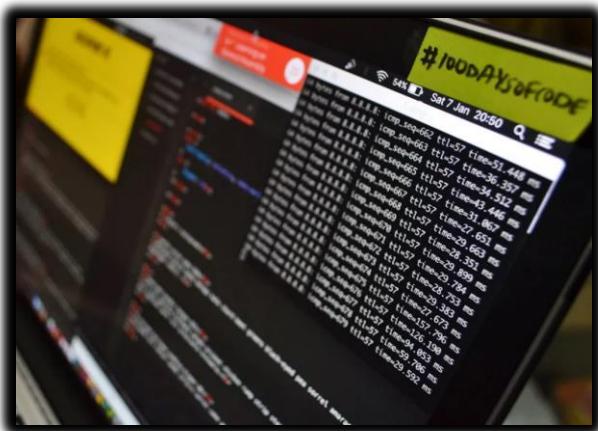


II. Reusable Code

Since the pointer can be changed at runtime, the **same instructions can operate on different data items**.

This makes indirect addressing ideal for:

- Loops
- Arrays
- Functions that process multiple variables



III. Examples of Indirect Addressing

Here are some examples in assembly language:

```
; Load the content of memory address stored at memory address 100 to the register R1.  
LOAD R1, @100  
  
; Load the content of the memory address stored at register R2 to register R1.  
LOAD R1, @R2  
  
; Store the content of register R1 to the memory address stored at register R3.  
STORE R1, @R3
```

```
int main() {  
    int array[10] = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10};  
    int *pointer = &array[10];  
  
    ;Load the element at the address pointed to by the pointer to the register EAX.  
    EAX = *pointer;  
  
    ;Return the value in register EAX.  
    return EAX;  
}
```

Indirect addressing is a total game-changer, but it's got some quirks you should know about.

One catch is that it's a bit slower than direct addressing.

This happens because the CPU can't just jump to the data—it has to take an extra step to look up where the pointer is actually aiming first.

Another thing is that it can be a bit risky. If your pointer isn't set up right, or if it's pointing to a spot in memory that doesn't exist, your program is going to crash or act weird.

As for the code:

```
int *pointer = &array[5];
```

Is it actually necessary to include that [5]? If you just write &array, you're grabbing the address of the very first item. By adding the [5], you're telling the computer,

"I don't want the start; I want to point specifically at the sixth item."

Technically, you don't *have* to use brackets to get there. These two lines do the exact same thing:

```
int *pointer = &array;  
int *pointer = &array[5];
```

The reason why you can omit the 5 in brackets is because arrays are implicitly converted to pointers to their first element.

This means that the expression `&array` is equivalent to the expression `&array[0]`.

However, there are some cases where it is good practice to explicitly specify the array index.

For example, if you are declaring a pointer to the last element of an array, it is more clear to write something like this:

```
int *pointer = &array[9];
```

Writing out the index makes it crystal clear that your pointer is aiming at the end of the array rather than the beginning. It's a "**tell, don't hide**" approach to coding.

It's also a smart move when you're working with functions or libraries. If a specific tool expects to start working from a certain spot in your list, being explicit prevents the program from getting lost.

Ultimately, it's your call whether to include the index or not. Just keep in mind that leaving it out changes where the pointer lands, so you have to be sure you're aiming at the right target.

Another Example:

Think about a string of text like "HELLO".

- If you use `&string`, you're pointing at the 'H'.
- If you use `&string[4]`, you're pointing specifically at the 'O'.

If you pass just `&string` to a function that's supposed to delete the last letter, the function has to do extra work to find the end. If you pass `&string[4]`, you've handed it the exact spot it needs to work on.

Another Example:

```
.data  
    byteVal BYTE 10h  
.code  
    mov esi, OFFSET byteVal  
    mov al, [esi]           ;al = 10h
```

Basically, you're just setting a marker in the computer's memory.

The first line creates a tiny piece of data (a byte) named **byteVal** and gives it the value 10h.

The second line takes the "home address" of that data and saves it into the **esi** register.

Think of the **offset** as a distance: it's the exact number of bytes you'd have to walk from the very start of the data section to reach the front door of **byteVal**.

Since **byteVal** is sitting in the data segment, its offset is just its position relative to the start of that "storage room."

By putting that address into **esi**, you're essentially giving the computer a pointer so it knows exactly where to look later.

The following assembly language instruction moves the offset of the variable **byteVal** into the register **esi**:

```
mov ESI, OFFSET byteVal
```

After this instruction is executed, the register **esi** will contain the offset of the variable **byteVal**.

The offset of a variable can be used to load or store the value of the variable.

For example, the following assembly language instruction loads the value of the variable **byteVal** into the register **al**:

```
mov al, [ESI]
```

To put it simply, an **offset** is just a GPS coordinate for your data.

When you run an instruction like **mov al, [byteVal]**, you're telling the CPU to go to a specific "distance" from the start of memory to find your value.

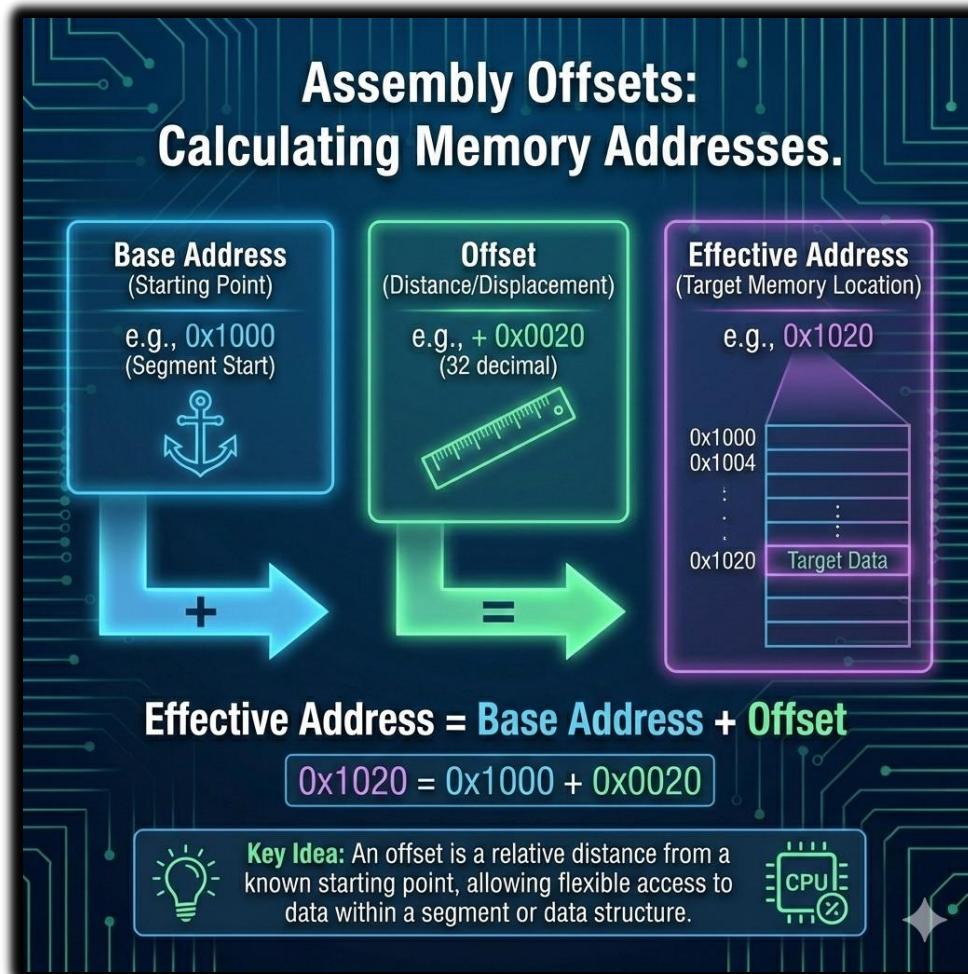
How the Compiler Uses Offsets for Functions

When your code calls a function, the compiler has to do some heavy lifting behind the scenes using these offsets:

- **Setting the Table:** Before the function starts, the compiler uses offsets to "push" your arguments onto the stack. It needs to know exactly where each piece of data sits so it can line them up correctly.
- **The Big Jump:** It uses the function's own offset to know where the actual instructions for that function are located in memory.
- **Cleaning Up:** Once the function is done, the compiler uses offsets again to "pop" those arguments off the stack and grab the return value.

The Bottom Line

Think of the **offset** as the "starting point" for everything.



Whether you're just moving a single byte into a register or handling a complex function call with multiple variables, the computer relies on these relative distances to make sure it's always grabbing the right data from the right drawer.



Using OFFSET to Access Memory

In this code, we have a variable **byteVal** declared in the .data section with the value 10h (16 in decimal).

The **OFFSET** keyword calculates the **memory address** of byteVal and stores it in the ESI register. Here's what's happening step by step:

1. byteVal is a **byte variable**, so it occupies **1 byte** in memory.
2. **OFFSET** byteVal tells the assembler to compute the **address of byteVal relative to the start of the data segment**.
3. Since byteVal is the **first item** in the data segment, its offset is 0x0.
4. `mov esi, OFFSET byteVal` loads the offset 0x0 into the ESI register.
5. The next instruction, `mov al, [esi]`, reads the byte **at the memory address stored in ESI** (which is the address of byteVal) and stores it in the AL register. Now, AL contains 10h (16 decimal).

Summary:

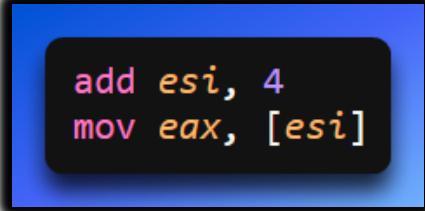
The **OFFSET** keyword gives the relative memory address of a variable in the data segment, allowing registers like ESI to act as **pointers** for accessing the variable.

Accessing Arrays with OFFSET

```
.data  
arrayD DWORD 10000h, 20000h, 30000h ; Define an array of doublewords  
  
.code  
mov esi, OFFSET arrayD ; Load the memory address of the array into ESI  
  
; Access the first number  
mov eax, [esi] ; Load the first doubleword into EAX  
; ESI points to the first doubleword  
  
; Access the second number  
add esi, 4 ; Move ESI to the next doubleword (increment by 4 bytes)  
mov eax, [esi] ; Load the second doubleword into EAX  
; ESI points to the second doubleword  
  
; Access the third number  
add esi, 4 ; Move ESI to the next doubleword (increment by 4 bytes)  
mov eax, [esi] ; Load the third doubleword into EAX  
; ESI points to the third doubleword
```

The same concept works for arrays. Suppose we have an array of three **doublewords** called arrayD:

1. The **OFFSET** keyword calculates the starting address of arrayD.
2. The ESI register is used as a pointer to traverse the array.
3. To move to the **next element**, you add 4 to ESI (because each doubleword is 4 bytes):



```
add esi, 4  
mov eax, [esi]
```

This allows you to **sequentially access each element** of the array in memory.

For example, if arrayD starts at **offset 10200h**, ESI will first point to 10200h, then 10204h, then 10208h, and so on, allowing each element to be loaded into EAX one by one.

Offset	Value
10200	10000h
10204	20000h
10208	30000h

← [esi]
← [esi] + 4
← [esi] + 8

INDEXED ADDRESSING

Indexed addressing is a way of finding the address of a value in memory. It works by **adding a constant (called a displacement) to the contents of a special register called an index register**.

- **Index register:** A CPU register that stores a base value used to offset an address.
- **Displacement:** A constant added to the index register to calculate the final memory address.

This mode is especially handy for working with arrays or other data structures like linked lists and trees, because it lets you **jump directly to a specific element** without knowing the absolute memory address.

I. Example: Accessing an array element

LOAD R1, [EAX + 4]

- R2 = base address of the array
- 5 = index/displacement
- This loads the 6th element (index 5) into register R1

The **effective address** of the element is calculated as:

```
Effective Address = Contents of R2 + 5
```

II. Example: Accessing a linked list element

```
LOAD R1, [EAX + 4]
```

- EAX = address of the current element
- 4 = offset to the next element
- Loads the next element into R1

Indexed addressing makes it easy to **navigate data structures**, but it has some limitations:

- Can be slower than direct addressing (CPU needs extra calculation).
- If the index register isn't initialized correctly, it can lead to errors.

Indexed Operands

Indexed operands are the combination of a **register (index)** and a **constant (displacement)** to access memory.

They're very useful for arrays because you don't need to know an element's exact memory address—just its index.

I. Examples in assembly:

```
LOAD R1, 5(R2)      ; Load array[5] into R1
LOAD R1, [EAX + 4]    ; Load next element in linked list into R1
STORE R1, 100(R4)     ; Store R1 at memory location R4 + 100
```

Equivalent C code:

```
int array[10] = {1,2,3,4,5,6,7,8,9,10};  
int index = 5;  
int element = array[index]; // Access array element using an index  
return element;
```

II. Indexed Operands with Scale Factors

When working with arrays where elements vary in size (like words, doublewords, etc.), **scale factors** are used to calculate the correct offset.

- **Scale factor:** The size of one array element in bytes
 - Word = 2 bytes
 - Doubleword = 4 bytes

III. Two common forms of indexed operands:

```
constant[reg]  
[constant + reg]
```

IV. Example: Byte array

```
.data  
arrayB BYTE 16h, 26h, 36h  
  
.code  
mov esi, 0 ; Initialize index register  
mov al, arrayB[esi] ; AL = 16h (first element)  
mov al, [arrayB + esi] ; Same result
```

V. Example: Word array with displacements

```
.data  
arrayW WORD 1090h, 2009h, 3EEEh  
  
.code  
mov esi, OFFSET arrayW  
mov ax, [esi]          ; AX = 1090h (first element)  
mov ax, [esi + 2]      ; AX = 2009h (second element)  
mov ax, [esi + 4]      ; AX = 3EEEh (third element)
```

VI. Example: Doubleword array with scale factor

```
.data  
arrayD DWORD 100h, 200h, 300h, 400h  
  
.code  
mov esi, 3 * TYPE arrayD ; Offset for arrayD[3] (3*4 = 12 bytes)  
mov eax, [arrayD + esi]   ; EAX = 400h
```

- `TYPE arrayD` gives the size of each element (4 bytes here).
- $3 * 4 = 12 \rightarrow$ the fourth element (index 3) is 12 bytes from the start.

Important tip: Arrays in assembly are zero-indexed:

- `arrayD[0]` → first element
- `arrayD[3]` → fourth element

So, setting `esi = 12` correctly accesses the element containing `400h`.

Quick Tips

1. Always **initialize the index register** before using it.
2. Use **displacements or scale factors** to navigate arrays of different sizes.
3. Double-check your offsets to avoid accessing the wrong memory location.
4. Indexed addressing is powerful but can be slower than direct addressing.

IMMEDIATE ADDRESSING MODE

Immediate addressing is the simplest way to provide data to an instruction.

Instead of going to memory or a register to fetch a value, you **put the constant value directly into the instruction**.

- Used for instructions like MOV, ADD, SUB, etc.
- The operand is **embedded in the instruction itself**, not stored elsewhere.

```
mov eax, 100
```

- Here, 100 is the immediate value.
- It's directly loaded into the EAX register.

Compare that with **direct addressing**:

```
mov ebx, [direct_value]
```

- direct_value is a memory location.
- The instruction fetches the value stored at that memory address and loads it into EBX.

So, the key difference:

Addressing Type	What Happens
Immediate	The value is in the instruction itself.
Direct	The value is in memory, and the instruction points to its address.

```

.data
    constant_value DWORD 42
    direct_value DWORD 10

.code
main PROC
    ; Immediate Addressing
    mov eax, 100      ; Load the immediate value 100 into EAX
    add eax, constant_value ; Add the constant value to EAX
    ; EAX now contains the result of 100 + 42 = 142

    ; Direct Addressing
    mov ebx, [direct_value] ; Load the value at the memory location pointed to by direct_value into EBX
    ; EBX now contains the value stored at the memory location pointed to by direct_value

    ; Exit the program
    invoke ExitProcess, 0
main ENDP
END main

```

Comparing the Five Common Addressing Modes

Here's a simple breakdown using a small assembly snippet:

```

mov eax, [array]      ; Direct addressing
mov ebx, [ptr]        ; Indirect addressing
mov ecx, [array + 4]  ; Indexed addressing
mov edx, 100          ; Immediate addressing
push 200              ; Stack addressing
pop stackVal         ; Stack addressing

```

What each does:

1. **Direct addressing:** `mov eax, [array]`
 - Loads the value at the memory location named array into EAX.
2. **Indirect addressing:** `mov ebx, [ptr]`
 - Loads the value from the memory location **pointed to by ptr** into EBX.
3. **Indexed addressing:** `mov ecx, [array + 4]`
 - Adds an offset (4 bytes here) to array and loads that element into ECX.
 - Great for accessing array elements.

4. **Immediate addressing:** mov edx, 100

- Loads the constant value 100 directly into EDX.

5. **Stack addressing:** push 200 / pop stackVal

- push 200 puts the value 200 on top of the stack.
- pop stackVal retrieves the top value from the stack into the variable stackVal.

```
.data
array DWORD 10, 20, 30, 40 ; An array for indexed addressing
value DWORD 100           ; A constant value for immediate addressing
ptr DWORD offset array    ; A pointer for indirect addressing
stackVal DWORD ?          ; A variable for stack addressing

.code
main PROC
; Direct Addressing
    mov eax, [array]    ; Load the value at the memory location pointed to by array into EAX

; Indirect Addressing
    mov ebx, [ptr]      ; Load the value at the memory location pointed to by ptr into EBX

; Indexed Addressing
    mov ecx, [array + 4] ; Load the value at the memory location (array + 4) into ECX
; This demonstrates indexed addressing by accessing the second element of the array

; Immediate Addressing
    mov edx, value       ; Load the immediate value 100 into EDX

; Stack Addressing
    push 200             ; Push the immediate value 200 onto the stack
    pop stackVal         ; Pop the value from the stack into stackVal

; Exit the program
    invoke ExitProcess, 0
main ENDP
END main
```

Quick Tip

- **Immediate = in the instruction**
- **Direct = value at memory location**
- **Indirect = memory address in a register**
- **Indexed = base address + offset**
- **Stack = value on the stack**

This table really helps visualize the differences:

MODE	OPERAND LOCATION	EXAMPLE
Immediate	Inside instruction	<code>mov eax, 100</code>
Direct	Fixed memory address	<code>mov eax, [array]</code>
Indirect	Address stored in register	<code>mov ebx, [ptr]</code>
Indexed	Base + offset	<code>mov ecx, [array + 4]</code>
Stack	Top of stack	<code>push 200 / pop</code>

Mode	How it works	Analogy
Immediate	Value is in the code	"Here is \$5."
Direct	Uses a variable name	"Go to the safe."
Indirect	Uses a pointer	"Go where the map says."
Indexed	Base address + Math	"Go 3 houses down."
Stack	Last-In, First-Out	"Stack of plates."

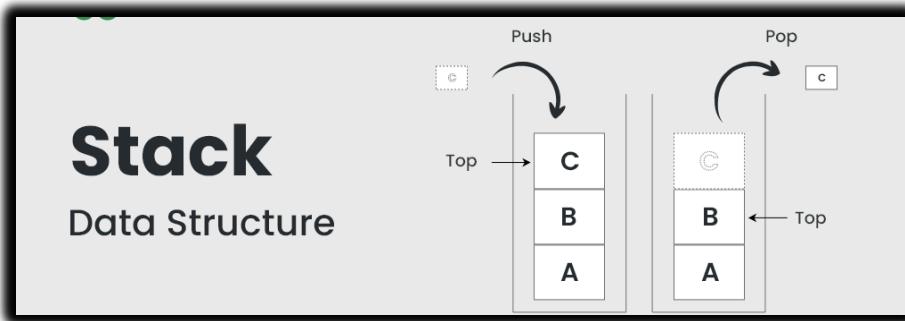
STACK ADDRESSING

Stack addressing is a key concept in computer architecture and assembly language, especially in stack-oriented architectures like x86.

It uses a **special memory region called the stack** to manage data and control program flow.

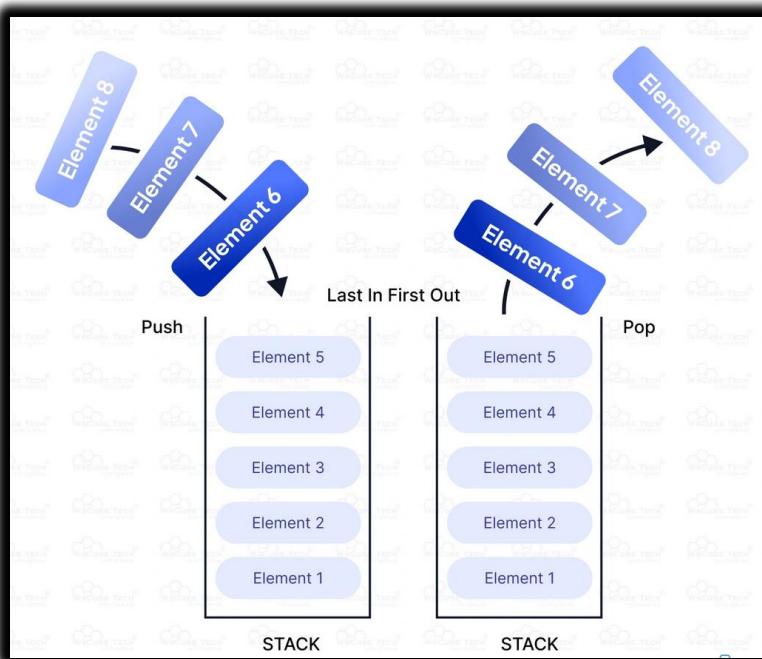
I. What is the Stack?

- The stack is a **temporary storage area in memory**.
- It works as a **Last-In, First-Out (LIFO)** structure: The **last item pushed** onto the stack is the **first one popped** off.



II. Stack Pointer (SP)

- The stack is managed with a special register called the **Stack Pointer (SP)**.
- The SP **points to the top of the stack**.
- **Operations:**
 - **Push:** Decrement SP, then store data at the address SP points to.
 - **Pop:** Retrieve data from the address SP points to, then increment SP.



III. Stack Operations

1. **Push:** Place a value on top of the stack.
 - Example: PUSH 200
 - SP moves down, and the value is stored at the new top.
2. **Pop:** Remove the top value from the stack.
 - Example: POP stackVal
 - The value is retrieved, and SP moves back up.

IV. Example in x86 Assembly

```
PUSH EAX ; Push the value in EAX onto the stack
POP EBX ; Pop the top value from the stack into EBX
```

PUSH adds a plate to the top, and **POP** takes the top one off. It's the primary way assembly handles function calls:

1. **The Call:** The computer pushes the **return address** (so it knows where to go back to) and any **parameters** the function needs onto the stack.
2. **The Work:** The function uses the stack to store its own temporary **local variables**.
3. **The Return:** Once finished, it pops those values off to clean up and uses the return address to jump back to the main code.

V. The "CALL" and "RET" Loop

```
; Call a function. Inside the function, it may push and pop values onto the stack
CALL MyFunction
; Return from the function, which pops the return address from the stack
RET
```

- **CALL:** Pushes the return address (the next instruction) onto the stack, then jumps to the function.
- **Function:** Can use the stack to save registers or store local variables. Remember to pop what you pushed before returning.
- **RET:** Pops the return address from the stack and jumps back to it.

TL;DR: CALL saves your place; RET returns to it.

VI. Passing Data via the Stack

```
;Push the argument to the function `add()` onto the stack.  
PUSH EAX  
  
;Call the function `add()`.  
CALL add  
  
;The return value of the function `add()` is now on the top of the stack.  
;Pop the return value off the stack and store it in the register `EAX`.  
POP EAX
```

1. **Push arguments:** Before calling a function, push the function's arguments onto the stack.
2. **Call the function:** CALL pushes the return address onto the stack and jumps to the function.
3. **Return value:** After the function finishes, the return value is pushed onto the stack.
4. **Pop result:** Retrieve the return value by popping it into a register.

Quick Tips for Stack Addressing:

- The stack is **Last-In, First-Out (LIFO)**: the last value pushed is the first one popped.
- The **stack pointer (SP)** keeps track of the top of the stack.
- Used for:
 - Function arguments
 - Return values
 - Temporary storage during operations

```
; Push the first operand onto the stack.  
PUSH EAX  
  
; Push the second operand onto the stack.  
PUSH EBX  
  
; Add the two operands on the stack and store the result in the register `EAX`.  
ADD EAX, [ESP] ; ESP is the stack pointer  
  
; Pop the second operand off the stack.  
POP EBX  
  
; Pop the first operand off the stack.  
POP EAX
```

Step-by-Step:

1. **PUSH EAX / PUSH EBX:** Save the values of EAX and EBX onto the stack.
2. **ADD EAX, [ESP]:** Add the value at the top of the stack (ESP points here) to EAX.
3. **POP EBX:** Restore the original value of EBX from the stack.
4. **POP EAX:** Retrieve the result of the addition back into EAX.

Notes:

- This is an **unconventional way** of doing addition using the stack.
- Typically, you would add registers directly (ADD EAX, EBX).
- The stack here is used to temporarily store operands and results.
- **Precondition:** EAX and EBX must already contain values before pushing.

```
.data
value1 DWORD 10    ; Define a 32-bit integer with the value 10
value2 DWORD 20    ; Define another 32-bit integer with the value 20

.code
main PROC
    ; Push the first value onto the stack
    PUSH DWORD PTR [value1]

    ; Push the second value onto the stack
    PUSH DWORD PTR [value2]

    ; Call a function to add the two values
    CALL add_values

    ; The result is now in EAX; you can use it or print it
    ; For demonstration purposes, we'll print the result
    MOV EAX, DWORD PTR [ESP]    ; Load the result from the stack to EAX
    CALL print_result

    ; Clean up the stack
    ADD ESP, 8    ; Remove the two values (4 bytes each)

    ; Exit the program
    MOV EAX, 1    ; Exit code 1
    INT 20h      ; DOS system call for program exit
main ENDP
```

```

; Function to add two values
add_values PROC
    POP EBX      ; Pop the second value into EBX
    POP EAX      ; Pop the first value into EAX
    ADD EAX, EBX ; Add the values, result in EAX
    PUSH EAX     ; Push the result back onto the stack
    RET
add_values ENDP

; Function to print the result (placeholder)
print_result PROC
    ; (Code for printing the value goes here - platform-dependent)
    RET
print_result ENDP

```

Data Section

- Defines two 32-bit integers: value1 = 10 and value2 = 20.

Code Section

- main PROC marks the program's entry point.

Stack Operations

- PUSH is used to place value1 and value2 onto the stack as 32-bit integers.
- CALL add_values calls a function to add the two numbers.
- After the function returns, the result is left on top of the stack.

Retrieving and Using the Result

- The result is moved from the stack into a register (e.g., EAX).
- Stack cleanup is done with ADD ESP, 8 to remove the two pushed values.

Function (add_values)

- POP retrieves value2 into EBX and value1 into EAX.
- ADD EAX, EBX computes the sum.
- The result is pushed back onto the stack for the caller.
- RET returns control to the calling code.

Program Exit

- Exit code is set with a register (e.g., EAX = 1) and a system call (e.g., INT 20h).

Summary

- This program demonstrates **stack-based addition**: push arguments, call a function, compute, push result, retrieve result, clean stack, and exit.
- Key stack concepts: LIFO order, using PUSH/POP, and function return with RET.

```
#include <stdio.h>

int addValues(int a, int b) {
    return a + b;
}

int main() {
    int value1 = 10;
    int value2 = 20;

    int result = addValues(value1, value2);

    printf("Result: %d\n", result);

    return 0;
}
```

The function addValues takes two integers, adds them, and returns the result.

This corresponds to the add_values function in assembly.

In main:

- value1 = 10 and value2 = 20 correspond to the assembly variables value1 and value2.
- addValues(value1, value2) performs the addition, just like the assembly function.
- The result is stored in result and printed with printf.

Summary:

- This C code performs the **same addition operation** as the MASM assembly code.
- The difference is that C uses **higher-level syntax** and **function calls** instead of manual stack operations.

POINTERS IN ASSEMBLY LANGUAGE

A **pointer** is a variable that stores the **memory address** of another variable. Pointers are essential in assembly for **direct memory manipulation**, efficient array access, and dynamic data handling.

Pointer Basics

- Pointers allow you to access and modify memory at runtime.
- In **32-bit mode**, pointers are typically **doublewords (DWORD)**.
- The PTR keyword tells the assembler the type of the pointer (BYTE PTR, WORD PTR, DWORD PTR).
- The OFFSET operator returns the address of a variable.

I. Example: Pointer to a Byte Array

```
.data
arrayB BYTE 19h, 29h, 39h, 49h      ; Define a byte array
ptrB DWORD OFFSET arrayB            ; Pointer to the start of arrayB
```

- ptrB stores the **memory address** of arrayB.
- Using OFFSET makes it clear that ptrB points to arrayB.

II. Near Pointers

- In 32-bit assembly, pointers are usually stored in **DWORDs** (4 bytes).
- Example with bytes and words:

```
.data
arrayB BYTE 19h, 29h, 39h, 49h
arrayW WORD 1999h, 2999h, 3999h

ptrB DWORD OFFSET arrayB ; Pointer to byte array
ptrW DWORD OFFSET arrayW ; Pointer to word array
```

III. Pointer Operations

Declaring and Using a Pointer

```
.data  
myByte BYTE 50h  
myPointer BYTE PTR ? ; Declare a pointer to a byte variable
```

Storing an Address

```
MOV myPointer, OFFSET myByte ; Store address of myByte in pointer
```

Accessing the Value

```
MOV EAX, [myPointer] ; Load the value pointed to by myPointer into EAX
```

Modifying the Value

```
INC [myPointer] ; Increment the value at the address stored in myPointer
```

Explanation:

1. Declare myPointer as a pointer to a byte.
2. Store the **address** of myByte in myPointer using OFFSET.
3. Load the **value** at that address into EAX by **dereferencing** the pointer ([myPointer]).
4. Increment the value at the pointer's address using INC [myPointer].

Why Pointers Are Useful

- Efficiently manipulate **arrays and data structures**.
- Dynamically allocate or modify memory at **runtime**.
- Pass arguments to functions **by reference**.
- Implement **function pointers** or jump tables.

TYPEDEF OPERATOR

The TYPEDEF operator allows you to create a user-defined type that has all the status of a built-in type when defining variables.

This is useful for creating pointer types, as it makes the code more readable and maintainable.

For example, the following declaration creates a new data type called PBYTE that is a pointer to bytes:

```
PBYTE TYPEDEF PTR BYTE
```

This declaration would usually be placed near the beginning of a program, before the data segment. Then, variables could be defined using the PBYTE type:

```
PBYTE TYPEDEF PTR BYTE

.data
    arrayB BYTE 10h,20h,30h,40h
    ptr1 PBYTE ? ; uninitialized
    ptr2 PBYTE arrayB ; points to an array
```

Advantages of using the TYPEDEF operator to create pointer types:

It makes the code more readable and maintainable, as it is clear what type of data the pointer is pointing to.

It can help to prevent errors, as the compiler will check that the pointer is only used to access the type of data that it is pointing to.

It can make the code more portable, as the same TYPEDEF declaration can be used on different platforms.

The following program demonstrates how to use the TYPEDEF operator to create pointer types:

```
.386
.model flat,stdcall
.stack 4096
ExitProcess proto,dwExitCode:dword

; Create user-defined types.
PBYTE TYPEDEF PTR BYTE ; pointer to bytes
WORD TYPEDEF PTR WORD ; pointer to words
DWORD TYPEDEF PTR DWORD ; pointer to doublewords

.data
    arrayB BYTE 10h,20h,30h,40h
    arrayW WORD 1,2,3
    arrayD DWORD 4,5,6

    ; Create some pointer variables.
    ptr1 PBYTE arrayB
    ptr2 WORD arrayW
    ptr3 DWORD arrayD

.code
main PROC

    ; Use the pointers to access data.
    mov esi, ptr1
    mov al, [esi] ; AL = 10h
    mov esi, ptr2
    mov ax, [esi] ; AX = 1
    mov esi, ptr3
    mov eax, [esi] ; EAX = 4

    invoke ExitProcess, 0

main ENDP
END main
```

The TYPEDEF operator is a powerful tool that can be used to create pointer types. This can make the code more readable, maintainable, portable, and error-prone.

JMP AND LOOP INSTRUCTIONS

The JMP and LOOP instructions in assembly language are used to implement conditional and unconditional transfers of control, respectively.

The **JMP instruction** causes an **unconditional transfer** to a destination, identified by a code label.

The **LOOP instruction** **repeats a block of statements** a specific number of times, using the **ECX register as a counter**.

Here is an example of how to use the JMP and LOOP instructions to implement a simple loop:

```
loop_start:  
    ;loop body  
    ...  
    ;check the loop condition  
    cmp ecx, 0  
    je loop_end  
  
    dec ecx ;decrement the loop counter  
    jmp loop_start ;jump to the beginning of the loop  
  
loop_end  
    ;loop's exit code
```

This loop will execute as long as the ECX register is not equal to zero. When ECX reaches zero, the loop will exit.

The LOOP instruction is a more concise and efficient way to implement a loop than using a JMP instruction.

The LOOP instruction automatically decrements the ECX register and performs the jump to the loop label, so you don't have to explicitly write these instructions yourself.

Here is an example of how to use the LOOP instruction to implement the same loop as above:

```
loop_start:  
    ; loop body  
    ...  
  
    ; loop condition  
    cmp ecx, 0  
    je loop_end  
  
    ; loop  
    loop loop_start  
  
loop_end:  
    ; loop exit code
```

This loop is equivalent to the previous loop, but it is more concise and efficient.

You can also use the LOOP instruction to implement more complex loops, such as loops that iterate over an array or loops that exit based on a different condition.

```
.data  
    count DWORD ?  
  
.code  
    mov ecx, 100 ; set loop count  
    top:  
        mov count, ecx ; save the count  
  
        ; modify ECX  
  
        mov ecx, count ; restore loop count  
    loop top  
  
    ; loop exit code
```

The loop count is initialized to 100 by moving the value 100 into the ECX register.

The loop counter is saved in the count variable by moving the value of the ECX register to the count variable.

The ECX register is modified. The loop counter is restored by moving the value of the count variable to the ECX register.

The LOOP instruction is executed, which decrements the ECX register and jumps to the top label if the ECX register is not equal to zero.

The loop exit code is executed if the ECX register is equal to zero. The following is a more detailed explanation of each step:

top:

loop top

This instruction decrements the ECX register and jumps to the top label if the ECX register is not equal to zero. This causes the loop to repeat until the ECX register is equal to zero.

Nested Loops

```
.data
    count DWORD ?
.code
    ; set outer loop count
    mov ecx, 100

L1:
    ; save outer loop count
    mov count, ecx

    ; set inner loop count
    mov ecx, 20

L2:
    ; loop body
    ...
    ; decrement inner loop counter
    dec ecx
    ; jump to beginning of inner loop if inner loop counter is not zero
    jne L2
    ; restore outer loop count
    mov ecx, count
    ; decrement outer loop counter
    dec ecx
    ; jump to beginning of outer loop if outer loop counter is not zero
    jne L1
```

The nested loop you provided is an example of a combination of an outer loop and an inner loop, a common construct in programming. Let me explain what this nested loop does:

Outer Loop (L1):

- The outer loop is controlled by the ecx register and is initialized with the value 100.
- Inside the outer loop: The current value of ecx (the outer loop counter) is saved into the count variable.
- The inner loop is initialized by setting ecx to 20.

Inner Loop (L2):

- The inner loop is also controlled by the ecx register and is initialized with the value 20. Inside the inner loop:
 - There is a placeholder comment for the loop body, where actual operations or instructions would be performed.
 - dec ecx is used to decrement the inner loop counter.
 - jne L2 is a conditional jump instruction. It checks if the inner loop counter (ecx) is not equal to zero. If it's not zero, the program jumps back to the beginning of the inner loop (L2) to continue iterating.

Outer Loop Continuation:

- After the inner loop completes (when ecx becomes zero), the program restores the original value of ecx from the count variable.
- dec ecx is used to decrement the outer loop counter.
- jne L1 is another conditional jump instruction. It checks if the outer loop counter (ecx) is not equal to zero. If it's not zero, the program jumps back to the beginning of the outer loop (L1) to continue iterating.

Summary:

- The inner loop runs **20 times per outer loop iteration**.
- The outer loop runs **100 times**, giving a total of **$100 \times 20 = 2000$ iterations**.
- The loop body executes repeatedly for each inner loop iteration.

Summing an Array in Assembly

```
;NB: Don't indent assembly code like me, I just want clarity
.386
.model flat,stdcall
.stack 4096
ExitProcess proto,dwExitCode:dword

.data
    intarray DWORD 10000h,20000h,30000h,40000h

.code
main PROC
    ; 1: EDI = address of intarray
    mov edi, OFFSET intarray
    ; 3: sum = 0
    mov eax, 0
    ; 2: initialize loop counter
    mov ecx, LENGTHOF intarray
L1:
    ; 5: add an integer
    add eax, [edi]
    ; 6: point to next element
    add edi, TYPE intarray
    ; 7: repeat until ECX = 0
    loop L1
    ; Exit with success code
    invoke ExitProcess, 0
main ENDP
END main
```

This example shows how to **sum the elements of an array** using registers and a simple loop in x86 assembly. Let's go through it step by step.

I. Setting Up the Array and Main Function

- The code first defines intArray, which is an array of **four 16-bit integers**.
- The program defines a main() function—this is the **entry point** where execution begins.

II. Initializing Registers for the Loop

Before starting the loop, we prepare a few CPU registers:

1. **EDI** – Holds the **address of the array**. We'll use it as a pointer to the current element in the array.
2. **EAX** – Initialized to 0. This register will **accumulate the sum** of the array elements.
3. **ECX** – Set to the **length of the array**. This acts as the **loop counter**, telling the CPU how many elements to process.

III. The Loop (Label: L1)

The loop is where the actual summing happens. Here's what each step does:

1. **Add current array element to EAX**
 - The CPU reads the value at the memory address in EDI and adds it to EAX.
 - EAX now contains the sum of all elements processed so far.
2. **Move to the next element**
 - Increment EDI by the **size of one array element** (2 bytes for a 16-bit integer).
 - This updates EDI so it points to the next element in the array.
3. **Loop control**
 - Decrement ECX automatically each time through the loop.
 - If ECX is not zero, jump back to the L1 label to process the next element.
 - This continues until all elements have been added.

IV. Exiting the Program

After the loop finishes:

- The program calls ExitProcess to **exit cleanly**, returning a success code.
- At this point, EAX holds the **total sum** of the array elements.

V. Key Takeaways

- **Registers as helpers:**
 - EDI → points to array elements
 - EAX → accumulates the sum
 - ECX → counts iterations
- **Looping in assembly:**
 - Use a label (L1) as a loop start point.
 - Use a conditional jump (loop or jnz) to repeat until the counter reaches zero.
- **Memory addressing:**
 - Incrementing EDI by the element size moves the pointer to the next array element.

COPYING STRINGS

Using a loop that copies a string, represented as an array of bytes with a null terminator value.

Indexed addressing works well for this type of operation because the same index register references both strings.

The target string must have enough available space to receive the copied characters, including the null byte at the end:

```

.386
.model flat, stdcall
.stack 4096
ExitProcess PROTO ,dwExitCode:DWORD

.data
    source BYTE      "This is the source string",0
    target BYTE      SIZEOF source DUP(0)

.code
main PROC
    mov esi, 0          ;index register
    mov ecx, SIZEOF source ;loop counter

L1:
    mov al, source[esi]      ;get a character from source
    mov target[esi], al       ;store it in the target
    inc esi                 ;move to the next character

invoke ExitProcess, 0
main ENDP
END main

```

When you move a 32-bit constant to a 64-bit register, the upper 32 bits (bits 32–63) of the destination are cleared (equal to zero):

```
mov rax,0FFFFFFFh ; rax = 00000000FFFFFFF
```

When you move a 16-bit constant or an 8-bit constant into a 64-bit register, the upper bits are also cleared:

```
mov rax,06666h ; clears bits 16-63
mov rax,055h   ; clears bits 8-63
```

When you move memory operands into 64-bit registers, however, the results are mixed. For example, moving a 32-bit memory operand into EAX (the lower half of RAX) causes the upper 32 bits in RAX to be cleared:

```
.data
    myDword DWORD    80000000h
.code
    mov rax, 0xFFFFFFFFFFFFFFFh
    mov eax, myDword           ;RAX = 0000000800000000
```

But when you move an 8-bit or a 16-bit memory operand into the lower bits of RAX, the highest bits in the destination register are not affected:

The following assembly code example shows how moving an 8-bit or 16-bit memory operand into the lower bits of RAX does not affect the highest bits in the destination register:

```
.data
myByte BYTE 55h
myWord WORD 6666h

.code
; Move the 16-bit value from myWord into the lower bits of RAX.
mov ax, myWord

; Display the value of RAX in hexadecimal format.
; Output: 0066 (only the lower 16 bits are displayed)
mov edx, 1 ; standard output
mov eax, RAX
syscall

; Move the 8-bit value from myByte into the lower bits of RAX.
mov al, myByte

; Display the value of RAX in hexadecimal format.
; Output: 55 (only the lower 8 bits are displayed)
mov edx, 1 ; standard output
mov eax, RAX
syscall
```

1. Moving Data into RAX

Think of **RAX** as a massive 64-bit storage unit. When you move a smaller piece of data into it, like **AX** (16-bit) or **AL** (8-bit), you are only touching the "bottom floor" of that unit.

- **mov ax, myWord**: Only the last 16 bits of RAX change. The other 48 bits stay exactly as they were.
- **mov al, myByte**: Only the very last 8 bits change.

2. The MOVSXD Trick (Sign Extension)

MOVSXD is the "Smart Copy" for moving a 32-bit value into a 64-bit register.

- It doesn't just copy the number; it looks at the **sign bit** (the very first bit of the 32-bit number).
- If the number is negative, it fills all the new empty space in the 64-bit register with **1s** to keep it negative.
- In our example, ebx starts with 9..., which in binary starts with a **1**. So, MOVSXD floods the rest of RAX with Fs (all ones), turning the whole register into FFFFFFFFFFFFFFh.

3. Why bother?

We do this because computers are picky about math. You can't easily add a 32-bit number to a 64-bit number.

MOVSXD "inflates" the smaller number so it fits the 64-bit math operations without changing its value.

Here is an example of how the MOVSXD instruction can be used to prepare data for further processing:

```
; EBX contains a 32-bit signed value.  
; RAX is a 64-bit register.  
  
movsx rax, ebx ; Sign extend the 32-bit value in EBX to 64 bits.  
  
; RAX now contains a 64-bit signed value.
```

This code can then be used to perform a mathematical operation on the value in RAX, such as:

```
; Add 10 to the value in RAX.  
add rax, 10
```

The **ADD instruction** expects a 64-bit operand, so the sign extended value in RAX is used in the operation.

The **MOVSXD instruction** is a powerful tool that can be used to prepare data for further processing in a variety of ways.

Instruction	Source Size	Dest Size	What happens to "Top" bits?
MOV (8/16 to 64)	8/16-bit	64-bit	Nothing (unchanged)
MOV (32 to 64)	32-bit	64-bit	Zero-extended (Upper 32 bits = 0)
MOVSXD	32-bit	64-bit	Sign-extended (Fills with Sign Bit)

The OFFSET operator generates a 64-bit address, which must be held by a 64-bit register or variable. In the following example, we use the RSI register:

```
.data
    myArray WORD 10,20,30,40
.code
    mov rsi,OFFSET myArray
```

Data Section Allocates memory for 4 words (8 bytes total).

Code Section:

- OFFSET myArray calculates the **memory address** of the array.
- mov rsi, OFFSET myArray stores this address in the **64-bit RSI register**.
- RSI now points to the start of myArray in memory.

4. Why This Matters

- Storing the memory address in a register allows you to **manipulate the array** efficiently.
- Enables **pointer-based operations**, indexed addressing, or iteration in loops.
- Important in **64-bit mode**, since registers like RSI, RDI, RCX, and RAX can hold full 64-bit addresses.

5. Additional Notes on 64-Bit Assembly

- **Loop Counters:** RCX is often used with the LOOP instruction.
- **64-bit integers and registers** simplify arithmetic and memory operations.
- **ASCII strings** (byte arrays) commonly use **indirect or indexed addressing** to access elements.

This sets the stage for our **64-bit version of SumArray**, which can iterate through the array using the RSI pointer.

```
; Summing an Array (SumArray_64.asm)
ExitProcess PROTO

.data
intarray QWORD 1000000000000h,2000000000000h,3000000000000h,4000000000000h

.code
main PROC
    mov rdi,OFFSET intarray      ; RDI = address of intarray
    mov rcx,LENGTHOF intarray   ; initialize loop counter
    mov rax,0                   ; sum = 0

L1: ; mark beginning of loop
    add rax,[rdi]                ; add a quadword
    add rdi,TYPE intarray        ; point to next element
    loop L1                      ; repeat until RCX = 0

    mov ecx,0                    ; ExitProcess return value
    call ExitProcess
main ENDP
END main
```

The **SumArray** program is a straightforward way to see how 64-bit assembly handles math on a larger scale.

Instead of the old 32-bit registers, this version uses the 64-bit heavyweights: **RAX** for the running total, **RDI** to point to the data, and **RCX** as the loop counter. It uses the **QWORD** directive to tell the assembler that every number in the array is a full 64 bits wide.

The logic is simple: the code walks through the array, adding each quadword into **RAX** one by one. The loop keeps spinning until **RCX** hits zero. Once the math is done, the program clears out **ECX** to signal a clean exit and calls **ExitProcess**.

6. Addition and Subtraction in 64-Bit Mode

In 64-bit mode, ADD and SUB work exactly like they do in 32-bit mode, just with more "elbow room."



You can manipulate massive numbers directly in the 64-bit registers without having to worry about breaking them into smaller chunks.

```
; Addition and Subtraction in 64-Bit Mode (AddSubtract_64.asm)

.data

.code
main PROC
    ; Adding 1 to a 32-bit number in RAX
    mov rax,0FFFFFFFh          ; fill the lower 32 bits
    add rax,1                  ; RAX = 10000000h

    ; Adding 16-bit values in AX and BX
    mov rax,0FFFFh            ; RAX = 000000000000FFFF
    mov bx,1                  ; BX = 0001
    add ax,bx                 ; RAX = 0000000000000000

    ; Subtracting 1 from zero in EAX
    mov rax,0                  ; RAX = 0000000000000000
    mov ebx,1                  ; EBX = 00000001
    sub eax,ebx                ; RAX = 00000000FFFFFF

    ; Subtracting 1 from zero in AX
    mov rax,0                  ; RAX = 0000000000000000
    mov bx,1                  ; BX = 0001
    sub ax,bx                 ; RAX = 000000000000FFFF

    ; Exit the program
    mov ecx,0                  ; ExitProcess return value
    call ExitProcess
main ENDP
END main
```

The ADD, SUB, INC, and DEC instructions in 64-bit mode affect CPU status flags similarly to 32-bit mode.

In an example, adding 1 to a 32-bit number in RAX causes bit 32 to receive a 1, resulting in RAX becoming 100000000h.

When using partial register operands, be aware that the remainder of the register is not modified. For example, adding 16-bit values in AX and BX does not affect the upper bits in RAX.

The same principle applies to subtraction. Subtracting 1 from zero in EAX or AX modifies the lower bits of RAX while leaving the upper bits unchanged.

These concepts outline the behavior of instructions and registers in 64-bit mode, particularly with regard to handling integers and loops.

Use of 64-bit Registers with Indirect Operands

Memory Increment and Decrement Instructions in 64-Bit Assembly

In 64-bit mode, when working with indirect operands, you must use a 64-bit general-purpose register to hold the memory address.

Additionally, it's essential to use the PTR operator to specify the size of the target operand. Here are some examples:

```
dec BYTE PTR [rdi]           ; Decrement an 8-bit target
inc WORD PTR [rbx]           ; Increment a 16-bit target
inc QWORD PTR [rsi]          ; Increment a 64-bit target
```

The following instructions demonstrate how to manipulate data directly in memory for different operand sizes using 64-bit registers:

I. dec BYTE PTR [rdi]

- **Instruction:** dec – decrements the operand by 1.
- **Operand size:** BYTE PTR indicates that the target is a single byte (8 bits).
- **Memory reference:** [rdi] points to the memory location addressed by the 64-bit register RDI.

Explanation:

This instruction **subtracts 1** from the byte stored at the memory location pointed to by RDI. It is commonly used when working with 8-bit data in memory, such as decrementing a counter or a byte variable.

II. inc WORD PTR [rbx]

- **Instruction:** inc – increments the operand by 1.
- **Operand size:** WORD PTR indicates that the target is a word (16 bits).
- **Memory reference:** [rbx] points to the memory location addressed by the 64-bit register RBX.

Explanation:

This instruction **adds 1** to the 16-bit value stored at the memory location pointed to by RBX. It is used for operations on 16-bit data, like incrementing a short integer in memory.

III. inc QWORD PTR [rsi]

- **Instruction:** inc – increments the operand by 1.
- **Operand size:** QWORD PTR indicates that the target is a quadword (64 bits).
- **Memory reference:** [rsi] points to the memory location addressed by the 64-bit register RSI.

Explanation:

This instruction **adds 1** to the 64-bit value stored at the memory location pointed to by RSI. It is suitable for operations involving 64-bit integers in memory, like counters or accumulators in 64-bit programs.

Summary

These instructions allow **direct memory manipulation** of values of different sizes:

Instruction	Size	Operation
dec BYTE PTR [rdi]	8 bits	Decrement by 1
inc WORD PTR [rbx]	16 bits	Increment by 1
inc QWORD PTR [rsi]	64 bits	Increment by 1

Scale Factors in Indexed Operands

In 64-bit mode, you can also use scale factors in indexed operands, similar to 32-bit mode.

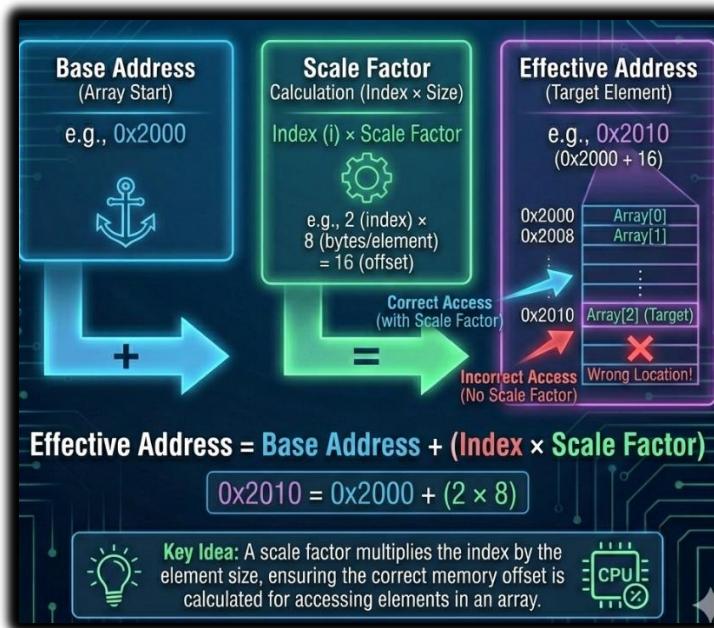
If you're working with an array of 64-bit integers, you can use a scale factor of 8 to correctly calculate offsets. Here's an example:

```
.data  
    array QWORD 1,2,3,4  
  
.code  
    mov esi,3           ; Subscript  
    mov eax, array[rsi*8] ; EAX = 4
```

In this example, we're accessing a 64-bit integer in the array using a scale factor of 8 to account for the size of each element.

A **scale factor** in assembly language is a multiplier used in **indexed addressing** to calculate the correct memory address of an element in an array. It accounts for the size of each element in the array.

For example, if you have an array of 64-bit integers (8 bytes each), using a scale factor of 8 ensures that when you access `array[i]`, the CPU correctly computes the memory offset as `base_address + i * 8`. Without the scale factor, the CPU would access the wrong memory location.



64-Bit Pointers in Assembly

In **64-bit mode**, pointers are **64-bit values**. That means they can hold the **memory address** of anything in your program—like pointing to a variable, an array, or any data in memory.

I. Example: Pointing to an Array

```
.data
    arrayB BYTE 10h,20h,30h,40h
    ptrB QWORD arrayB
```

- arrayB is just a simple array of 4 bytes: 10h, 20h, 30h, 40h.
- ptrB is a **64-bit variable** (QWORD) that **holds the address of arrayB**.

Think of it like this:

- arrayB = a row of boxes with numbers in them: [10][20][30][40]
- ptrB = a sticky note that says: “Start looking at box number X”
- The CPU doesn’t copy the array into ptrB; it just **stores the location** of the array.

II. Making It Clearer with OFFSET

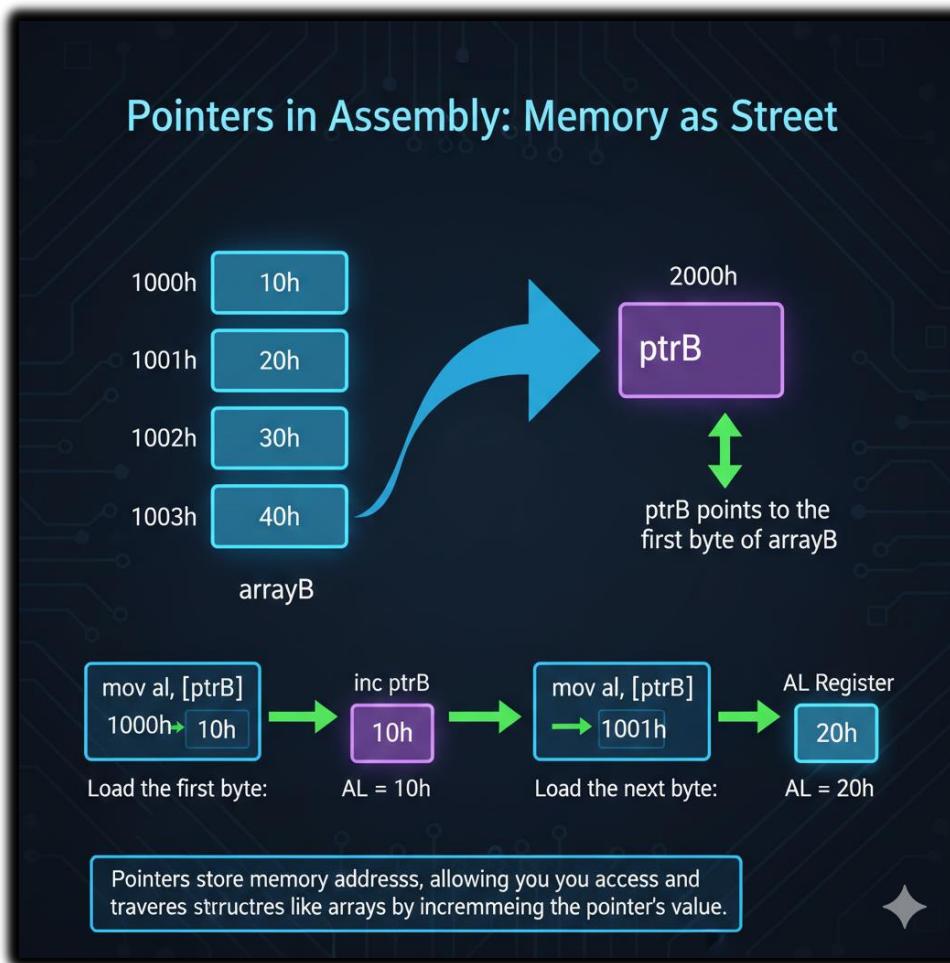
You can make the pointer’s relationship to the array more explicit using the **OFFSET** operator:

```
.data
    arrayB BYTE 10h,20h,30h,40h
    ptrB QWORD OFFSET arrayB
```

- OFFSET arrayB literally says: “**Store the address of arrayB in ptrB**”
- Functionally, it’s the same as before, but it’s clearer when reading the code.
- Beginners often think ptrB stores the array itself—**it does not!** It just points to where the array lives in memory.

III. Visualization

Imagine memory like a street with houses:



Imagine memory like a street with houses:

1000h → 10h

1001h → 20h

1002h → 30h

1003h → 40h

ptrB (at 2000h) = 1000h

ptrB doesn't copy the contents of arrayB.

It points to the first byte of arrayB.

If you want to read the array using ptrB, you tell the CPU:

```
mov al, [ptrB] ; Load the first byte of arrayB into AL
```

Now AL = 10h.

Increment the pointer to get the next byte:

```
inc ptrB  
mov al, [ptrB] ; Load the second byte (20h)
```

IV. Key Takeaways for Beginners

1. **64-bit pointers = 64-bit memory addresses**
2. A pointer **never stores the data itself**, only the location of the data
3. Using OFFSET makes the pointer's target explicit and easier to understand
4. You can then use the pointer to **read/write array elements** in memory

V. Beginner Analogy

- arrayB = row of numbered mailboxes
- ptrB = address written on a sticky note pointing to the first mailbox
- Reading/writing via ptrB = going to the mailbox the note points to and checking the content

CHAPTER SUMMARY: WHAT WE COVERED

In this chapter, we've covered various essential concepts related to data transfers, addressing, arithmetic, and loops in assembly language programming. Here's a concise summary of the key points:

MOV Instruction:

- The MOV instruction is used for transferring data from a source operand to a destination operand.

Zero and Sign Extension:

- MOVZX extends a smaller operand to a larger one by zero-filling.
- MOVSX extends a smaller operand to a larger one by sign extension.

XCHG Instruction:

- XCHG swaps the contents of two operands, with at least one of them being a register.

Operand Types:

- Direct operand represents the address of a variable.
- Direct-offset operand adds a displacement to a variable's name, generating a new offset.
- Indirect operand uses a register containing the address of data.
- Indexed operand combines a constant with an indirect operand for effective memory access.

Arithmetic Instructions:

- INC adds 1 to an operand.
- DEC subtracts 1 from an operand.
- ADD adds a source operand to a destination operand.
- SUB subtracts a source operand from a destination operand.
- NEG reverses the sign of an operand.

Status Flags:

- Various CPU status flags are affected by arithmetic operations, including Sign, Carry, Parity, Auxiliary Carry, Zero, and Overflow flags.

Operators:

- OFFSET operator returns the distance of a variable from the beginning of its segment.
- PTR operator overrides a variable's declared size.
- TYPE operator returns the size of a variable or array element.
- LENGTHOF operator returns the number of elements in an array.
- SIZEOF operator returns the number of bytes used by an array initializer.
- TYPEDEF operator creates user-defined types.

Loops:

- JMP instruction unconditionally branches to another location.
- LOOP instruction is used in counting-type loops, with ECX as the counter in 32-bit mode and RCX in 64-bit mode.

64-Bit Mode:

- In 64-bit mode, a 64-bit general-purpose register must be used when working with indirect operands.
- Scale factors can be used in indexed operands, useful for arrays of 64-bit integers.

QUESTIONS

Question 9: What value will RAX contain after the following instruction executes?
mov rax,44445555h

Answer: RAX will contain the hexadecimal value 44445555h.

Question 10: What value will RAX contain after the following instructions execute?

```
.data
    dwordVal DWORD 84326732h
.code
    mov rax,0FFFFFFF00000000h
    mov rax,dwordVal
```

Answer: RAX will contain the value FFFFFFFF84326732h.

Question 11: What value will EAX contain after the following instructions execute?

```
.data  
    dVal DWORD 12345678h  
.code  
    mov ax,3  
    mov WORD PTR dVal+2,ax  
    mov eax,dVal
```

Answer: EAX will contain the value 00030000h.

Question 12.0: What will EAX contain after the following instructions execute?

```
.data  
    dVal DWORD ?  
.code  
    mov dVal,12345678h  
    mov ax,WORD PTR dVal+2  
    add ax,3  
    mov WORD PTR dVal,ax  
    mov eax,dVal
```

Answer: EAX will contain the value 00030000h.

Question 12.5: Is it possible to set the Overflow flag if you add a positive integer to a negative integer?

Answer: Yes, it is possible to set the Overflow flag when adding a positive integer to a negative integer if the result is out of range for the destination operand.

Question 13: Will the Overflow flag be set if you add a negative integer to a negative integer and produce a positive result?

Answer: No, the Overflow flag will not be set when adding two negative integers that result in a positive value.

Question 14: Is it possible for the NEG instruction to set the Overflow flag?

Answer: Yes, the NEG instruction can set the Overflow flag if the operation results in an overflow.

Question 15: Is it possible for both the Sign and Zero flags to be set at the same time?

Answer: No, it is not possible for both the Sign and Zero flags to be set simultaneously. If the result is zero, the Zero flag is set, indicating that the sign is not negative, so the Sign flag is not set. If the result is non-zero, the Sign flag is set to indicate whether it's positive or negative.

Exercise 1: Converting from Big Endian to Little Endian

```
.data
bigEndian BYTE 12h,34h,56h,78h
littleEndian DWORD 0

.code
mov eax, 0
mov al, [bigEndian]
mov ah, [bigEndian + 1]
mov [littleEndian], ax

mov eax, 0
mov al, [bigEndian + 2]
mov ah, [bigEndian + 3]
mov [littleEndian + 2], ax
```

Indented:

```
.data
    bigEndian BYTE 12h,34h,56h,78h
    littleEndian DWORD 0

.code
    mov eax, 0
    mov al, [bigEndian]
    mov ah, [bigEndian + 1]
    mov [littleEndian], ax

    mov eax, 0
    mov al, [bigEndian + 2]
    mov ah, [bigEndian + 3]
    mov [littleEndian + 2], ax
```

This program copies the value from bigEndian to littleEndian while reversing the order of the bytes.

Exercise 2: Exchanging Pairs of Array Values

```
.data
    array DWORD 1, 2, 3, 4, 5, 6, 7, 8
    arraySize DWORD 8

.code
    mov esi, 0
    mov ecx, [arraySize]

    exchangeLoop:
        mov eax, [array + esi]      ; Load the current element
        add esi, 4                  ; Move to the next element
        xchg eax, [array + esi]    ; Exchange the values
        loop exchangeLoop          ; Repeat for all pairs
```

This program exchanges every pair of values in the array with an even number of elements.

Exercise 3: Summing the Gaps between Array Values

```
.data
    array DWORD 0, 2, 5, 9, 10
    arraySize DWORD 5
    sumOfGaps DWORD 0

.code
    mov esi, 0
    mov ecx, [arraySize] - 1      ; One less than the array size

    sumLoop:
        mov eax, [array + esi]    ; Load the current element
        add esi, 4                ; Move to the next element
        sub eax, [array + esi]    ; Calculate the gap
        add [sumOfGaps], eax      ; Add gap to the sum
        loop sumLoop              ; Repeat for all gaps
```

This program calculates the sum of all the gaps between successive array elements in the array variable.

You can assemble and run these programs using an x86 assembler like NASM or MASM in either 32-bit or 64-bit mode as per your preference.